

Relational Model and Relational Algebra

Dr Na Yao

Objectives

- Terminology of relational model.
- How tables are used to represent data.
- Properties of database relations.
- How to identify candidate, primary, alternate, and foreign keys.
- Meaning of entity integrity and referential integrity.
- How to form queries in relational algebra.

Relational model

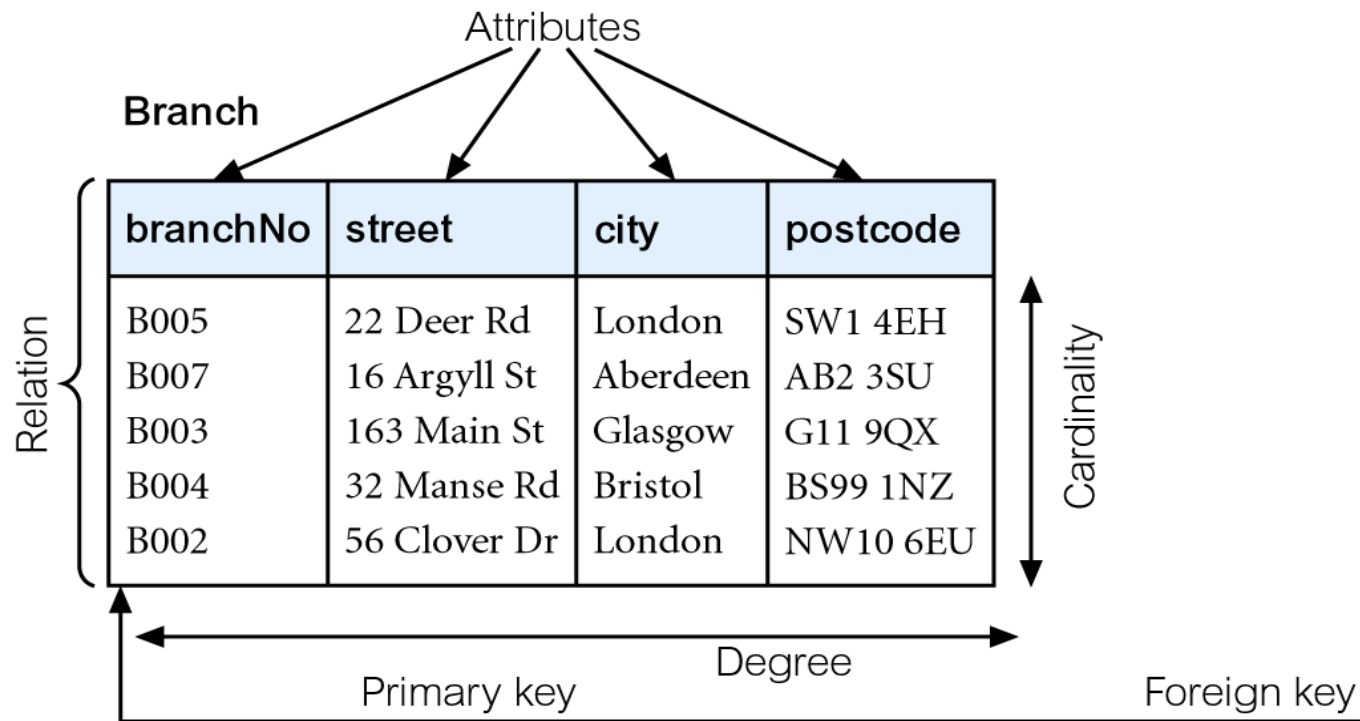
- E. F. Codd proposed relational data model in 1970.
- In the relational model, all data is logically structured within **relations** (tables).
- Each **relation** is made up of **attributes** (columns) of data.
- Each **tuple** (row) contains one value per attribute.

Terminology

- Relation: A relation is a table with columns and rows.
- Attribute: An attribute is a named column of a relation.
- Domain: the set of allowable values for one or more attributes.

Terminology

- Tuple: A tuple is row of a relation.
- Degree: the number of attributes in a relation.
- Cardinality: the number of tuples in a relation.
- Relational database: A collection of normalized relations with distinct relation names.



Staff

staffNo	fName	lName	position	sex	DOB	salary	branchNo
SL21	John	White	Manager	M	1-Oct-45	30000	B005
SG37	Ann	Beech	Assistant	F	10-Nov-60	12000	B003
SG14	David	Ford	Supervisor	M	24-Mar-58	18000	B003
SA9	Mary	Howe	Assistant	F	19-Feb-70	9000	B007
SG5	Susan	Brand	Manager	F	3-Jun-40	24000	B003
SL41	Julie	Lee	Assistant	F	13-Jun-65	9000	B005

Relation

Examples of Attribute Domains

Attribute	Domain Name	Meaning	Domain Definition
branchNo	BranchNumbers	The set of all possible branch numbers	character: size 4, range B001–B999
street	StreetNames	The set of all street names in Britain	character: size 25
city	CityNames	The set of all city names in Britain	character: size 15
postcode	Postcodes	The set of all postcodes in Britain	character: size 8
sex	Sex	The sex of a person	character: size 1, value M or F
DOB	DatesOfBirth	Possible values of staff birth dates	date, range from 1-Jan-20, format dd-mmm-yy
salary	Salaries	Possible values of staff salaries	monetary: 7 digits, range 6000.00–40000.00

Creating relations (tables) in SQL

```
Create Table Branch(branchNo, street, city,  
    postcode)
```

```
Create Table Branch(branchNo integer, street  
    string, city string, postcode char(6))
```


Properties of relations

- Relation name is distinct from all other relation names in relational schema.
- Each cell of relation contains exactly one atomic (single) value. (*First normal form*)
- Each attribute has a distinct name.
- Values of an attribute are all from the same domain.
- Each tuple is distinct; there are no duplicate tuples.
- Order of attributes has no significance.
- Order of tuples has no significance, theoretically.

Relational keys

- Candidate Key
 - A set of attributes that uniquely identifies a tuple within a relation.
 - Uniqueness : In each tuple, candidate key uniquely identify that tuple.
 - Irreducibility: No proper subset of the candidate key has the uniqueness property.
- Primary Key
 - Candidate key selected to identify tuples uniquely within relation.
- Foreign Key
 - Attribute, or set of attributes, within one relation that matches candidate key of some (possibly same) relation.

Representing Relational Schema

- Relation name (attribute 1, attribute 2, ... attribute n)
 - Branch (branchNo, street, city, postcode)
 - Staff (staffNo, fName, lName, position, sex, DOB, salary, branchNo)
- Exercise: do the relational schema for all the relations in *DreamHome* database.

Integrity Constraints

- Null
 - Represents value for an attribute that is currently unknown or not applicable for tuple.
 - Deals with incomplete or exceptional data.
 - Represents the absence of a value and is not the same as zero or spaces, which are values.

Integrity Constraints

- Entity Integrity

In a base relation, no attribute of a primary key can be null.

- Referential Integrity

If foreign key exists in a relation, either foreign key value must match a candidate key value of some tuple in its home relation or foreign key value must be null.

- General Constraints

Glossary


- A “crowd funded” glossary – every student can add entry and comment on an entry.

Students with good contribution will be rewarded 😊



Relational Algebra

- Relational algebra is formal language associated with the relational model.
- Relational algebra operations work on one or more relations to define another relation without changing the original relations.
- Allows expressions to be nested, just as in arithmetic.
This property is called closure.



“Thames Barrier Closure” by Chris Wheal, flickr

Relational Algebra

- Basic operations:
 - Selection
 - Projection
 - Cartesian product
 - Union
 - Set difference
- Additional operations:
 - Join
 - Intersection
 - Division

Selection

- $\sigma_{\text{condition}} (R)$

Works on a single relation R and defines a relation that contains only those tuples (rows) of R that satisfy the specified condition (predicate).

Example


- List all staff with a salary greater than £10,000.

$$\sigma_{\text{salary} > 10000} (\text{Staff})$$

- List all female staff with a salary greater than £10,000

$$\sigma_{\text{salary} > 10000 \wedge \text{sex} = \text{"F"}} (\text{Staff})$$

- List all branches in London.



Selection


Projection

- $\Pi_{\text{col1}, \dots, \text{coln}}(R)$

Works on a single relation R and defines a relation that contains a vertical subset of R, extracting the values of specified attributes and eliminating duplicates.

Example

- Produce a list of salaries for all staff, showing only staffNo, fName, lName, and salary details.



Projection

$$\Pi_{\text{staffNo, fName, lName, salary}}(\text{Staff})$$

Combine selection and projection

Example

- List staff number, position and salary of any staff whose salary is greater than £20,000.

Answer 1: $\sigma_{\text{salary} > 20000}(\Pi_{\text{staffNo}, \text{position}, \text{salary}}(\text{Staff}))$

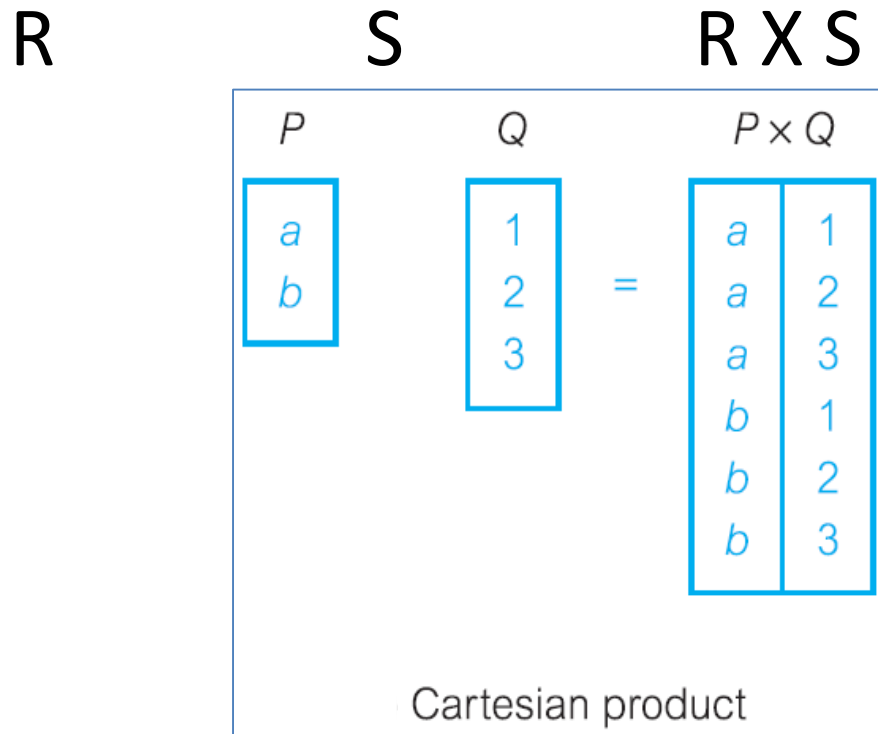
Answer 2: $\Pi_{\text{staffNo}, \text{position}, \text{salary}}(\sigma_{\text{salary} > 20000}(\text{Staff}))$

- List the branchNo of all branches in London.

Cartesian product

- $R \times S$

Defines a relation that is the concatenation of every tuple of relation R with every tuple of relation S.



Cartesian product

- $R \times S$

Defines a relation that is the concatenation of every tuple of relation R with every tuple of relation S.

R

S

$R \times S$

Example

- List the names and comments of all clients who have viewed a property for rent.

$(\Pi_{\text{clientNo}, \text{fName}, \text{lName}}(\text{Client})) \times (\Pi_{\text{clientNo}, \text{propertyNo}, \text{comment}}(\text{Viewing}))$

$\sigma_{\text{Client.clientNo} = \text{Viewing.clientNo}}(((\Pi_{\text{clientNo}, \text{fName}, \text{lName}}(\text{Client})) \times (\Pi_{\text{clientNo}, \text{propertyNo}, \text{comment}}(\text{Viewing}))))$

Join

- Derivative of Cartesian product, equivalent to performing a Selection operation, using the join condition as the selection formula, over the Cartesian product of the two operand relations.
- Natural Join
- Theta join

Natural Join

- $R \bowtie S$

- The result of the natural join is the set of all combinations of tuples in R and S that are equal on their common attribute names.

T		U		$T \bowtie U$		
A	B	B	C	A	B	C
a	1	1	x	a	1	x
b	2	1	y	a	1	y
		3	z			

Natural join

Natural Join

- $R \bowtie S$
 - The result of the natural join is the set of all combinations of tuples in R and S that are equal on their common attribute names.

Example

- List the names and comments of all clients who have viewed a property for rent.

$(\Pi_{\text{clientNo}, \text{fName}, \text{lName}}(\text{Client}) \bowtie$

$(\Pi_{\text{clientNo}, \text{propertyNo}, \text{comment}}(\text{Viewing}))$

$\Pi_{\text{clientNo}, \text{fName}, \text{lName}, \text{propertyNo}, \text{comment}} ((\text{Client}) \bowtie (\text{Viewing}))$

Theta join

- $R \bowtie_F S$
 - Defines a relation that contains tuples satisfying the condition F from the Cartesian product of R and S .
- Can rewrite Theta join using basic Selection and Cartesian product operations.

$$R \bowtie_F S = \sigma_F (R \times S)$$


Union

- $R \cup S$
 - Union of two relations R and S defines a relation that contains all the tuples of R, or S, or both R and S, duplicate tuples being eliminated.
 - R and S must be union-compatible.
- If R and S have I and J tuples, respectively, union is obtained by concatenating them into one relation with a maximum of $(I + J)$ tuples.

Example

- List all cities where there is either a branch office or a property for rent.

$$\Pi_{\text{city}}(\text{Branch}) \cup \Pi_{\text{city}}(\text{PropertyForRent})$$




Set difference

- $R - S$
 - Defines a relation consisting of the tuples that are in relation R , but not in S .
 - R and S must be union-compatible.

example

- List all cities where there is a branch office but no properties for rent.

$$\Pi_{\text{city}}(\text{Branch}) - \Pi_{\text{city}}(\text{PropertyForRent})$$




Intersection

- $R \cap S$
 - Defines a relation consisting of the set of all tuples that are in both R and S.
 - R and S must be union-compatible.
- Expressed using basic operations:
$$R \cap S = R - (R - S)$$

Example


- List all cities where there is both a branch office and at least one property for rent.

$$\Pi_{\text{city}}(\text{Branch}) \cap \Pi_{\text{city}}(\text{PropertyForRent})$$



Division

- Assume relation R is defined over the attribute set A and relation S is defined over the attribute set B such that $B \subseteq A$ (B is a subset of A). Let $C = A - B$, that is, C is the set of attributes of R that are not attributes of S .
- $R \div S$ defines a relation over the attributes C that consists of set of tuples from R that match combination of *every* tuple in S .



Division (shaded area)

Example of division

Example - Division

- Identify all clients who have viewed all properties with three rooms.

$$(\Pi_{\text{clientNo}, \text{propertyNo}}(\text{Viewing})) \div (\Pi_{\text{propertyNo}}(\sigma_{\text{rooms} = 3}(\text{PropertyForRent})))$$

$\Pi_{\text{clientNo}, \text{propertyNo}}(\text{Viewing})$

clientNo	propertyNo
CR56	PA14
CR76	PG4
CR56	PG4
CR62	PA14
CR56	PG36

$\Pi_{\text{propertyNo}}(\sigma_{\text{rooms} = 3}(\text{PropertyForRent}))$

propertyNo
PG4
PG36

RESULT

clientNo
CR56

Relational Algebra Exercise

Given the following relational schema:

Hotel(hotelNo, hotelName, city)

Room(roomNo, hotelNo, type, price)

Booking(hotelNo, guestNo, dateFrom, dateTo, roomNo)

Guest(guestNo, guestName, guestAddress)

Formulate the following queries in relational algebra:

1. Give the hotel number of those hotels with a room price greater than £50.
2. List the price and types of all rooms in Grosvenor Hotel.
3. Produce a relation containing all rooms at all hotels.
4. Give all hotel names with a room price above £50.
5. List the guest number of all guests currently staying at the Grosvenor Hotel. (Hint: use `current_date()` for today's date)
6. List guest number of all guests who have booked all hotels in Beijing.
7. List names of all guests who have booked all hotels in London.