

# DATABASE

## week 3 Summary

Dr Na Yao

So far

- 3.1 Normalization
- 3.2 Advanced Normalization
- 3.3 Transaction management
- 3.3 Distributed DBMS

**All Very  
important!!!**

# Normalization & Advanced Normalization

Everything is important, but  
this is the **MOST** important

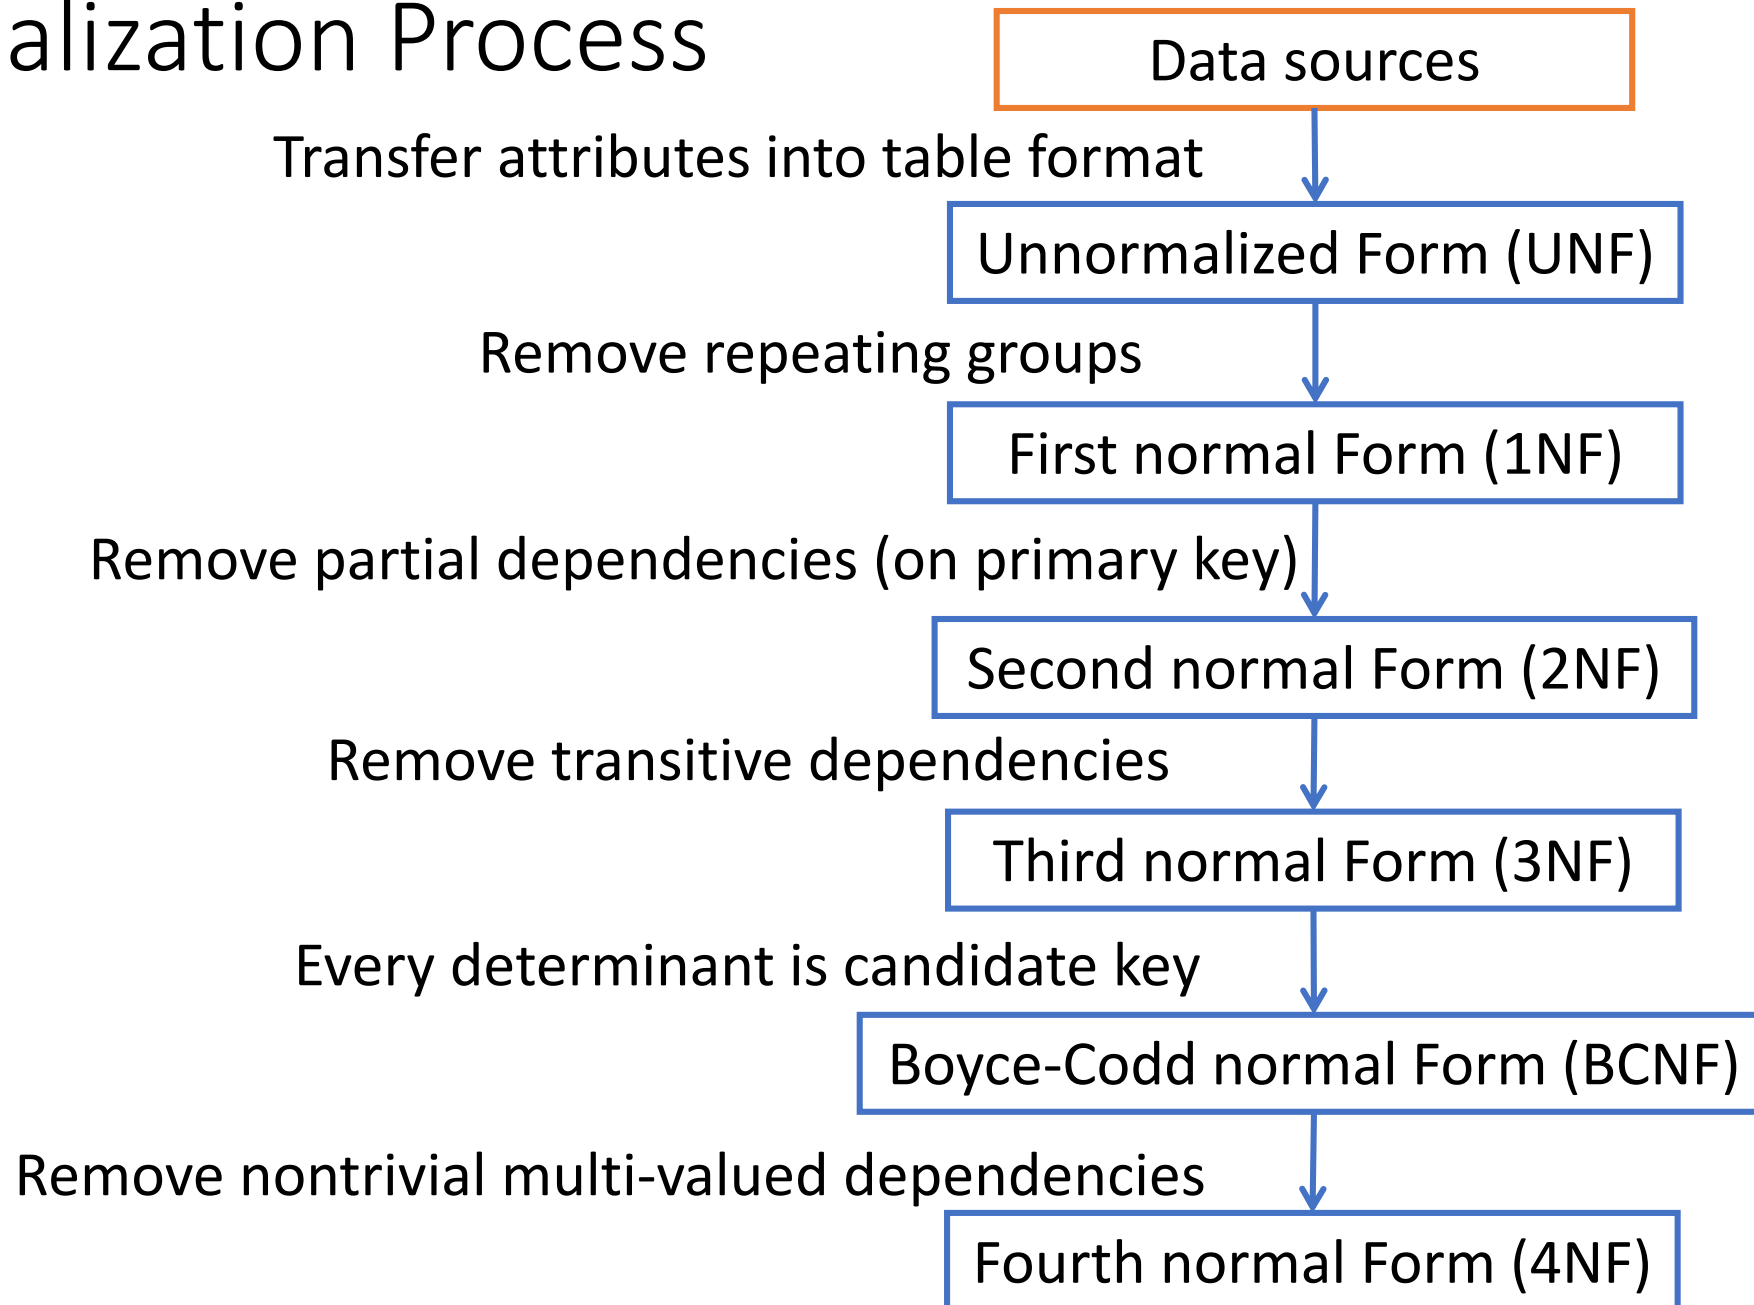
# Functional dependency

- Important concept used for normalization.
- Functional dependency describes relationship between **attributes**.
- For example, if A and B are attributes of relation R, B is functionally dependent on A (denoted  $A \rightarrow B$ ), if each value of A in R is associated with exactly one value of B in R.

# Other concepts

- Full functional dependency
- Partial functional dependency
- Transitive functional dependency
  
- Multi-valued dependency
- Trivial and nontrivial multi-valued dependency

# Normalization Process



Transaction management

Everything is important, but  
this is the **MOST** important

# ACID Properties of Transactions

- Atomicity
- Consistency
- Isolation
- Durability

I ❤️  
ACID



**CONCURRENCY =  
MANY TRANSACTIONS ARE SIMULTANEOUSLY RUN IN THE DATABASE**

# Concurrency problems

- Might **violate ACID** properties!!!!
  - different transactions may read/write the same data **concurrently**
  - breaking the **isolation** property.
- **3 main problems**
  1. Lost updates problem
  2. Uncommitted updates problem
  3. Incorrect analysis problem
- Solution: **2 Phase-Locking**




# Lost Update Problem

- Successfully completed update is **overridden** by another user.

Time	T1	T2	X
t <sub>1</sub>		begin_transaction	100
t <sub>2</sub>	begin_transaction	read(X)	100
t <sub>3</sub>	read(X)	X = X+100	100
t <sub>4</sub>	X = X-10	write(X)	200
t <sub>5</sub>	write(X)	commit	90
t <sub>6</sub>	commit		90


**Solution:** preventing T1 from reading X until after the update.



# Uncommitted Dependency Problem

Time	T3	T4	X
t <sub>1</sub>		begin_transaction	100
t <sub>2</sub>		read(X)	100
t <sub>3</sub>		X = X+100	100
t <sub>4</sub>	begin_transaction	write(X)	200
t <sub>5</sub>	read(X)	...	200
t <sub>6</sub>	X = X-10	rollback	100
t <sub>7</sub>	write(X)		190
t <sub>8</sub>	commit		190


- **Solution:** preventing T3 from reading X until after T4 commits or aborts.



# Inconsistent Analysis Problem *dirty/unrepeatable read*

Time	T5	T6	X	Y	Z	sum
t <sub>1</sub>		begin_transaction	100	50	25	
t <sub>2</sub>	begin_transaction	Sum = 0	100	50	25	0
t <sub>3</sub>	read(X)	read(X)	100	50	25	0
t <sub>4</sub>	X = X-10	sum = sum + X	100	50	25	100
t <sub>5</sub>	write(X)	read(Y)	90	50	25	100
t <sub>6</sub>	read(Z)	sum = sum + Y	90	50	25	150
t <sub>7</sub>	Z = Z + 10		90	50	25	150
t <sub>8</sub>	write(Z)		90	50	35	150
t <sub>9</sub>	commit	read(Z)	90	50	35	150
t <sub>10</sub>		sum = sum + Z	90	50	35	185
t <sub>11</sub>		commit	90	50	35	185

**Solution:** preventing T6 from reading X (and Z) until after T5 completed updates.



# Two-Phase Locking (2PL)



- All locking operations precede unlock operation in the transaction.
- Principle of 2PL
  - Every transaction must lock an item (read or write) before accessing it
  - Once a lock has been released, no new items can be locked



- Two phases for transaction:
  1. Growing phase - acquires all locks but cannot release any locks = **LOCKS**
  2. Shrinking phase - releases locks but cannot acquire any new locks = **UNLOCKS**

# Preventing the Lost Update Problem using 2PL

Time	T1	T2	X
$t_1$		begin_transaction	100
$t_2$	begin_transaction	write_lock(X)	100
$t_3$	write_lock(X)	read(X)	100
$t_4$	WAIT	$X := X + 100$	100
$t_5$	WAIT	write(X)	200
$t_6$	WAIT	commit/unlock(X)	200
$t_7$	read(X)		200
$t_8$	$X := X - 10$		200
$t_9$	write(X)		190
$t_{10}$	commit/unlock(X)		190

# Preventing the Uncommitted Dependency Problem using 2PL

Time	T3	T4	X
$t_1$		begin_transaction	100
$t_3$		write_lock( $X$ )	100
$t_4$		read( $X$ )	100
$t_5$	begin_transaction	$X := X + 100$	100
$t_4$	write_lock( $X$ )	write( $X$ )	200
$t_5$	WAIT	rollback/unlock( $X$ )	100
$t_4$	read( $X$ )		100
$t_6$	$X := X - 10$		100
$t_7$	write( $X$ )		90
$t_8$	commit/unlock( $X$ )		90



Time	T5	T6	X	Y	Z	sum
t <sub>1</sub>		begin_transaction	100	50	25	
t <sub>2</sub>	begin_transaction	Sum = 0	100	50	25	0
t <sub>3</sub>	write_lock(X)		100	50	25	0
t <sub>4</sub>	read(X)	read_lock(X)	100	50	25	0
t <sub>5</sub>	X = X-10	WAIT	100	50	25	0
t <sub>6</sub>	write(X)	WAIT	90	50	25	0
t <sub>7</sub>	write_lock(Z)	WAIT	90	50	25	0
t <sub>8</sub>	read(Z)	WAIT	90	50	25	0
t <sub>9</sub>	Z = Z + 10	WAIT	90	50	25	0
t <sub>10</sub>	write(Z)	WAIT	90	50	35	0
t <sub>11</sub>	Commit/unlock(X,Z)	WAIT	90	50	35	0
t <sub>12</sub>		read(X)	90	50	35	0
t <sub>13</sub>		sum = sum + X	90	50	35	90
t <sub>14</sub>		read_lock(Y)	90	50	35	90
t <sub>15</sub>		read(Y)	90	50	35	90
t <sub>16</sub>		sum = sum + Y	90	50	35	140
t <sub>17</sub>		read_lock(Z)	90	50	35	140
t <sub>18</sub>		read(Z)	90	50	35	140
t <sub>19</sub>		sum = sum + Z	90	50	35	175
t <sub>20</sub>		commit	90	50	35	175

# DDBMS

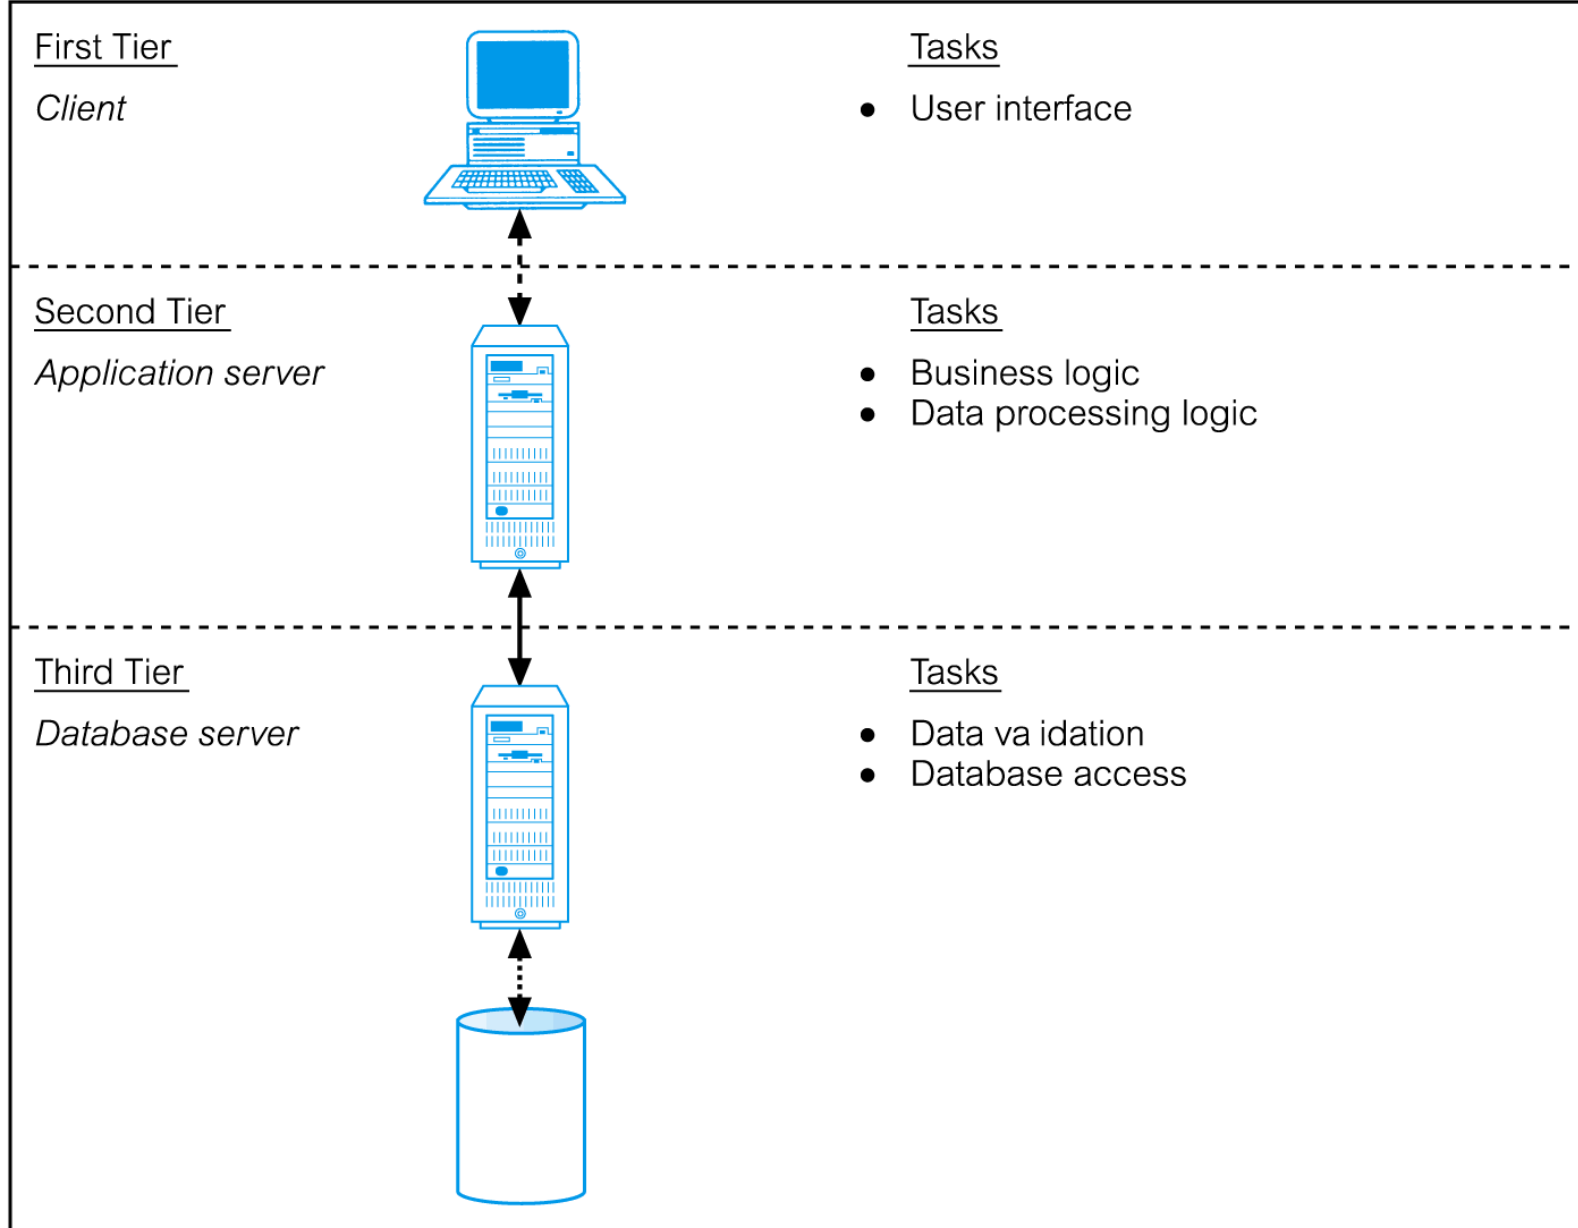
Everything is important, but  
this is the **MOST** important

# Client/Server Architecture

1. **The user interface layer:** runs on the end-user's computer (the *client*).
2. **The business logic and data processing layer:** this middle tier runs on a server and is often called the *application server*.
3. **A DBMS:** stores the data required by the middle tier. This tier may run on a separate server called the *database server*.

(The three-tier architecture can be extended to *multi*-tiers, with additional tiers added to provide more flexibility and scalability.)

# Three-Tier Client-Server Architecture




# Distributed Processing $\neq$ Distributed DBMS


- Distributed Processing:
  - centralized database
  - that can be accessed over a computer network



- Distributed DBMS:
  - Single logical database that is split into fragments
  - These fragments are distributed over a computer network



**DDBMS**



**Distributed Processing**

# DDBMS Design Methodology

- **Normalization** to produce a design for the global relations.
- Examine **topology** of system to determine where databases will be located.
- Analyse most important transactions and identify appropriateness of horizontal/vertical **fragmentation**.
- Decide which relations are **not to be fragmented**.
- Examine **relations** on 1 side of relationships and determine a suitable fragmentation schema. Relations on many sides may be suitable for horizontal derived fragmentation.
- Check for **situations where either vertical or mixed fragmentation would be appropriate** (that is, where transactions require access to a subset of the attributes of a relation).

# DDBMS vs Relational DBMS Design Methodology

- **Normalization** to produce a design for the global relations.
- Examine **topology** of system to determine where databases will be located.
- Analyse most important transactions and identify appropriateness of horizontal/vertical **fragmentation**.
- Decide which relations are **not to be fragmented**.
- Examine **relations** on 1 side of relationships and determine a suitable fragmentation schema. Relations on many sides may be suitable for horizontal derived fragmentation.
- Check for **situations where either vertical or mixed fragmentation would be appropriate** (that is, where transactions require access to a subset of the attributes of a relation).





**Understand IT ALL!!!!!**