

NoSQL

Dr Na Yao

Objectives

- Understand the motivation of NoSQL
- Understand and be able to explain the concepts of NoSQL
- Understand and be able to explain the application areas of NoSQL

Big data

- In the past decade the amount of data being created has skyrocketed.
 - The rate of data creation is only accelerating.
 - The data we deal with is diverse.
-
- Traditional database systems, such as relational databases, have been pushed to the limit, and failed to scale to “**Big Data**”.
 - To tackle the challenges of Big Data, a new breed of technologies has emerged. Many of these new technologies have been grouped under the term **NoSQL**.

HOW TO WRITE A CV



The NoSQL Movement

- An emerging “movement” around non-relational software for Big Data
- Origin:
 - Google (BigTable, MapReduce framework)
 - Amazon (distributed key/value store called Dynamo).
- Vibrant Open Source community
 - followed with Hadoop, HBase, MongoDB, Cassandra, RabbitMQ, and countless other projects.
- Currently defined as “**what it’s not**”

NoSQL

- Not Only SQL or “Not Relational”.
- key features:
 - Flexible schema
 - Quicker/cheaper to set up
 - Massive scalability
 - Relaxed consistency --> higher performance & availability
- Disadvantages:
 - No declarative query language --> more programming
 - Relaxed consistency --> fewer guarantees

NoSQL

- NoSQL systems frequently do not provide ACID properties
- Updates are eventually propagated, but there are limited guarantees on the consistency of reads.
- “BASE” instead of “ACID”:
 - BASE = Basically Available, Soft state, Eventually consistent
 - ACID = Atomicity, Consistency, Isolation, Durability
- By giving up ACID constraints, one can achieve much higher performance and scalability.

Why NoSQL?


No ACID constraints

Gain performance and scalability

NoSQL systems


- Several incarnations
 - MapReduce framework
 - Key-value stores
 - Document stores
 - Graph database systems

MapReduce Framework



- Originated at Google
- Open source implementation: Hadoop
- For processing and generating big data sets with a parallel, distributed algorithm on a cluster.

MapReduce, Characteristics



- No data model, data stored in files
- Functions:
 - map(), filtering and sorting ← divide problems into subproblems
 - E.g., sorting students by first name into queues, one queue for each name.
 - reduce(), summary operation ← do work on subproblems + combine results
 - E.g., counting the number of students in each queue, yielding name frequencies.
- System provides data processing “glue”, fault-tolerance, scalability

Key-Value Stores

- Extremely simple interface:
 - Data model: (key, value) pairs
 - Operations: Insert(key,value), Fetch(key), Update(key), Delete(key)
- Implementation: efficiency, scalability, fault-tolerance
 - Records distributed to nodes based on key
 - Replication
 - Single-record transactions, “eventual consistency”
- Example systems
- Google BigTable, Amazon Dynamo, Cassandra, Voldemort, HBase, ...

Document Stores

- Like Key-Value Stores except value is document
 - Data model: (key, document) pairs
 - Document: JSON, XML, other semi-structured formats
 - Basic operations: Insert(key,document), Fetch(key),
Update(key), Delete(key)
 - Also Fetch based on document contents
- Example systems
 - CouchDB, MongoDB, SimpleDB, ...

Graph Database Systems

- Data model: nodes and edges
- Nodes may have properties (including ID)
- Edges may have labels or roles
- Interfaces and query languages vary
- Single-step versus “path expressions” versus full recursion
- Example systems
 - Neo4j, FlockDB, Pregel, ...

Why such (semi-structured) data stores are needed?

- Semi-structured or flat files based data stores are best for massive data that is read, possibly frequently, but with minimal updates
- There is much less overhead to process data in this format
- We also have the flexibility to process data that doesn't have a completely fixed structure

What NoSQL should **NOT** be used for

- Anything that requires **frequent updates** as well as reads, or that requires high integrity and atomicity (ACID properties)
- Examples are similar to transaction databases for inventory and financial records
- This is not just a question of massive data or distributed processing!
- There are large, distributed relational databases like Visa or Amazon that need more structured data with transaction semantics
- These applications are better suited to relational databases even at large scale

In a nutshell

- NoSQL: alternative, non-traditional DB technology to be used in large scale environments where (ACID) transactions are not a priority