

HADOOP ARCHITECTURE

CLOUD COMPUTING

Dr. Atm Shafiul Alam

a.alam@qmul.ac.uk

Queen Mary University of London

School of Electronic Engineering and Computer Science


QUICK RECAP...

Yesterday we...

- Learnt about the basics of distributed computing
 - Split data up & process in parallel
- Learnt about the principles of Map Reduce
 - A particular paradigm for “Big Data”
- Ran some through basic pseudocode

What is

So I will use a “Cloud Computing” platform to hire access to machines



Cloud Computing platforms allow you to hire machines on-demand to run your tasks!

What is cloud computing?

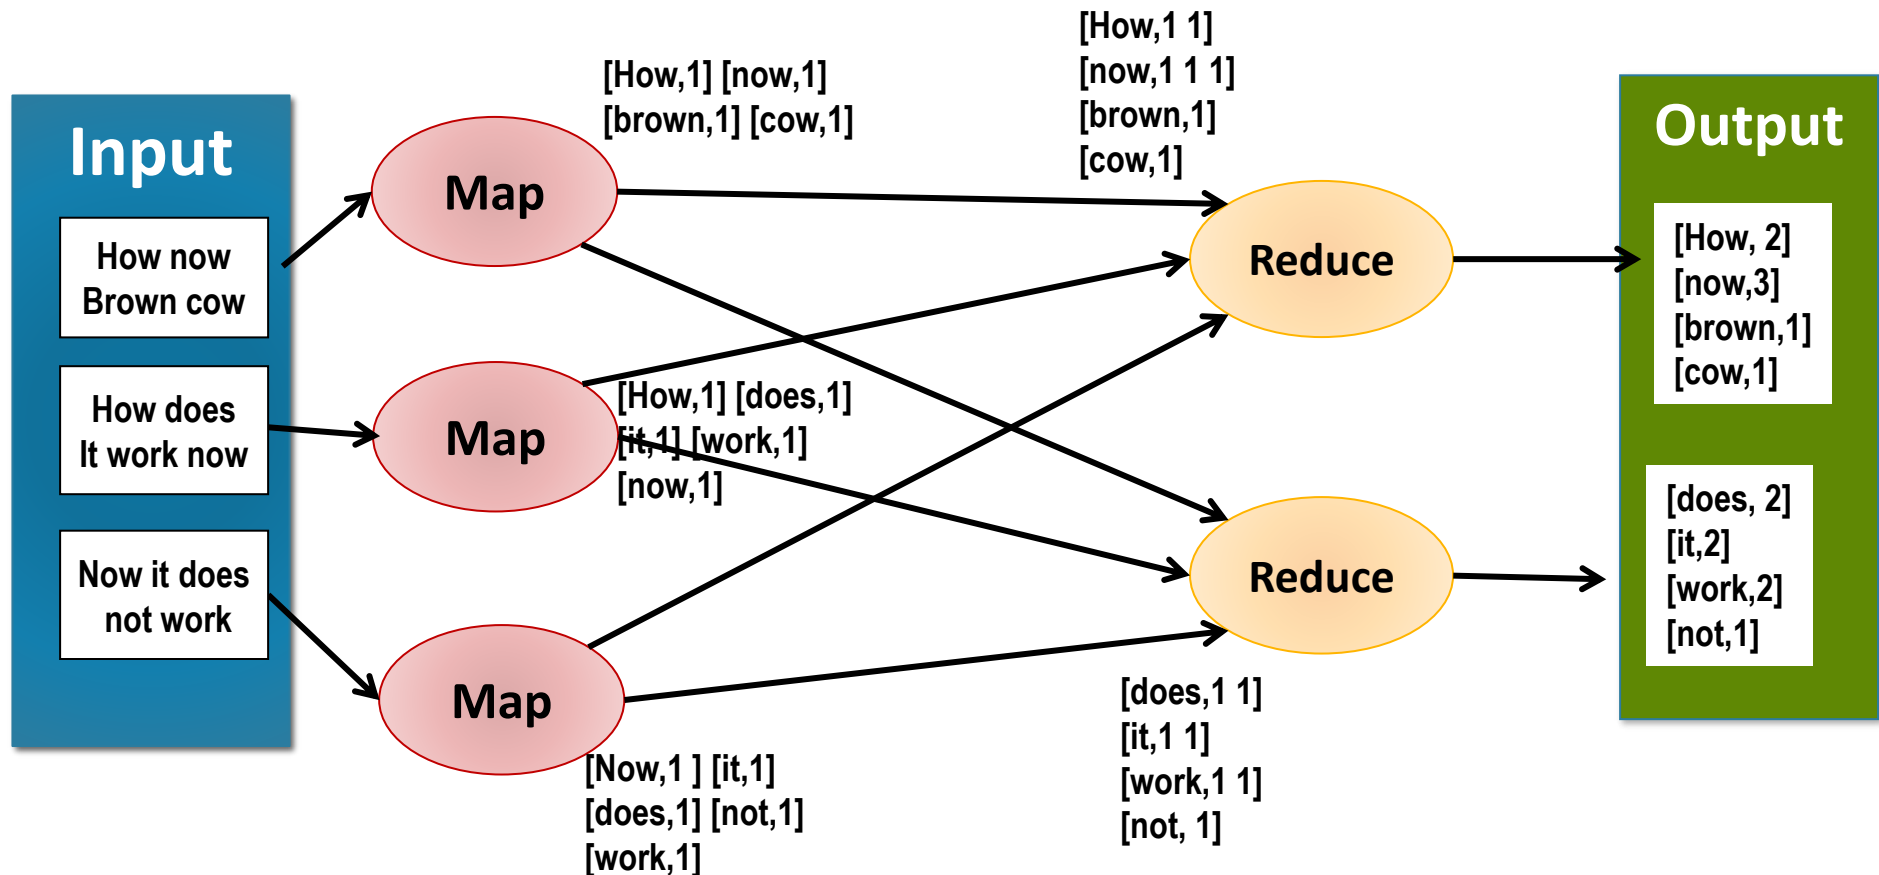
- You have already learnt more about how cloud computing works in Week 1.
- From now, we can simply assume that we have lots of computers to run our MapReduce application on
- ...but please feel free to ask if you're confused!

Parallel computing

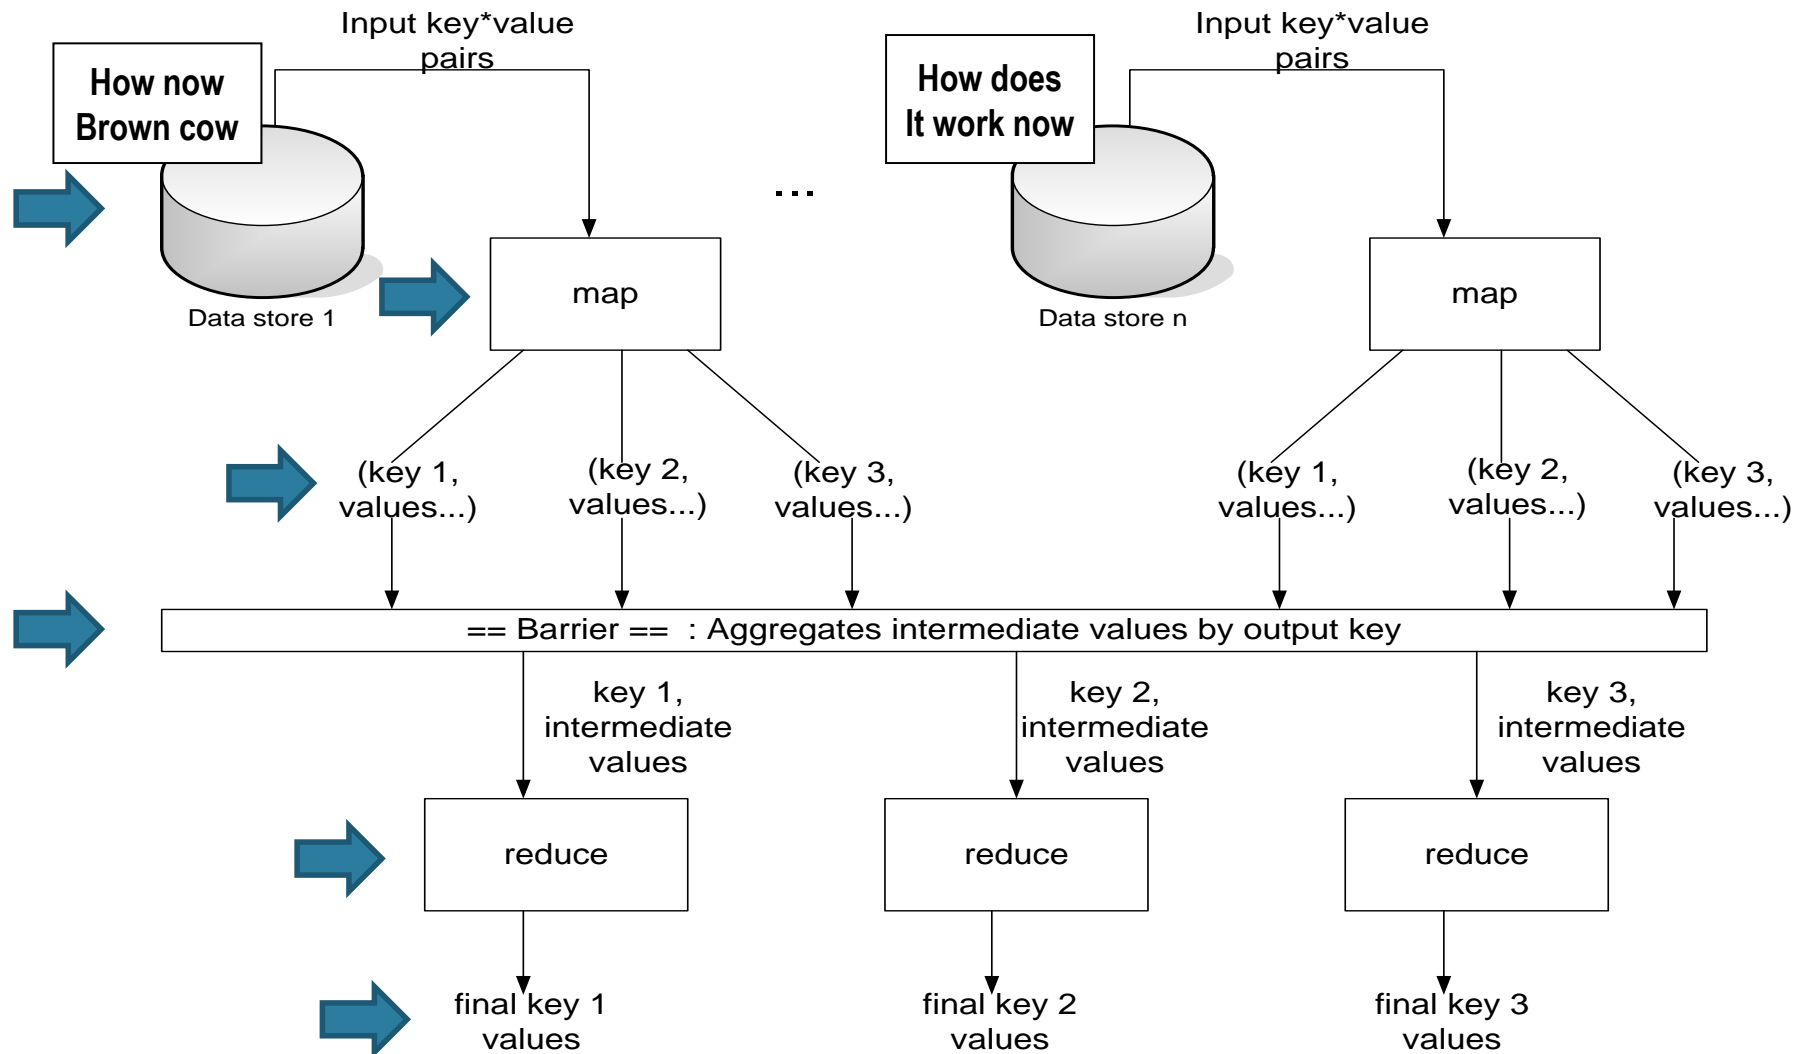
- The use of a number of **processors** to perform a calculation
- The calculation will be divided into tasks, sent to different processors (e.g. compute servers).
- Generally, we do this by **splitting data** across multiple **servers**

Note: cloud computing platforms are often used for parallel computing (because you can cheaply hire many computers)...

Word Count Example



Synchronisation and message passing



Today's contents

- **Anatomy of a MapReduce job**
- The Combiner
- Apache Hadoop
- Hadoop job execution: YARN
- Hadoop storage: HDFS

The Apache Hadoop project

- The brainchild of Doug Cutting (Yahoo)
- Apache open source framework (written in Java)
- Started in 2007 when code was spun out of Nutch
- Has grown into a large top-level project at Apache with significant ecosystem
 - V2 (YARN, 2013) structures it as a generic platform
 - Even third-party distros *a la* Linux (Cloudera, Hortonworks)



Hadoop Physical Requirements

- Designed to run in **clusters** of commodity PCs
 - Leverages heterogeneous capabilities
- Scales up to thousands of connected machines
- Suitable for Local Networks / Data Centers
 - Rack servers connected over a LAN
 - Clusters distributed over the Internet are not feasible
 - Network would become an enormous bottleneck (imagine sending terabytes over data over your DSL connection)

Hadoop Cluster



Principles of Hadoop Design

- *Data is distributed* around network
 - No centralized data server
 - Every node in cluster can host data
 - Data is replicated to support fault tolerance
- *Computation is sent to data*, rather than vice-versa
 - Code to be run is sent to nodes
 - Results of computations are aggregated as tend
- *Basic architecture is master/worker*
 - Master, *aka* JobNode, launches application
 - Workers, *aka* WorkerNodes, perform bulk of computation

Hadoop offers

- Redundant, Fault-tolerant data storage
- Parallel computation framework
- Job coordination



Hadoop offers

- Redundant, Fault-tolerant data storage
- Parallel computation framework
- Job coordination



Programmers

*No longer need to
worry about*



Q: Where file is located?

Q: How to handle failures &
data lost?

Q: How to divide computation?

Q: How to program for scaling?



Some MapReduce Terminology

- **Job** – A “full program” - an execution of a Mapper and Reducer across a dataset
- **Task** – An execution of a Mapper or a Reducer on a slice of data
 - a.k.a. Task-In-Progress (TIP)
- **Task Attempt** – A particular instance of an attempt to execute a task on a machine


Hadoop job

- A **Hadoop Job** is packaged as a **Jar** file containing all the code for Mapper and Reducer functions
- The job is assigned a cluster-unique ID
 - A set number of reattempts is managed for job tasks
- The file is replicated over the Hadoop nodes
 - Move computation to the data

Map/Reduce job



Map/Reduce framework roles




Job execution: Complete MapReduce job flow

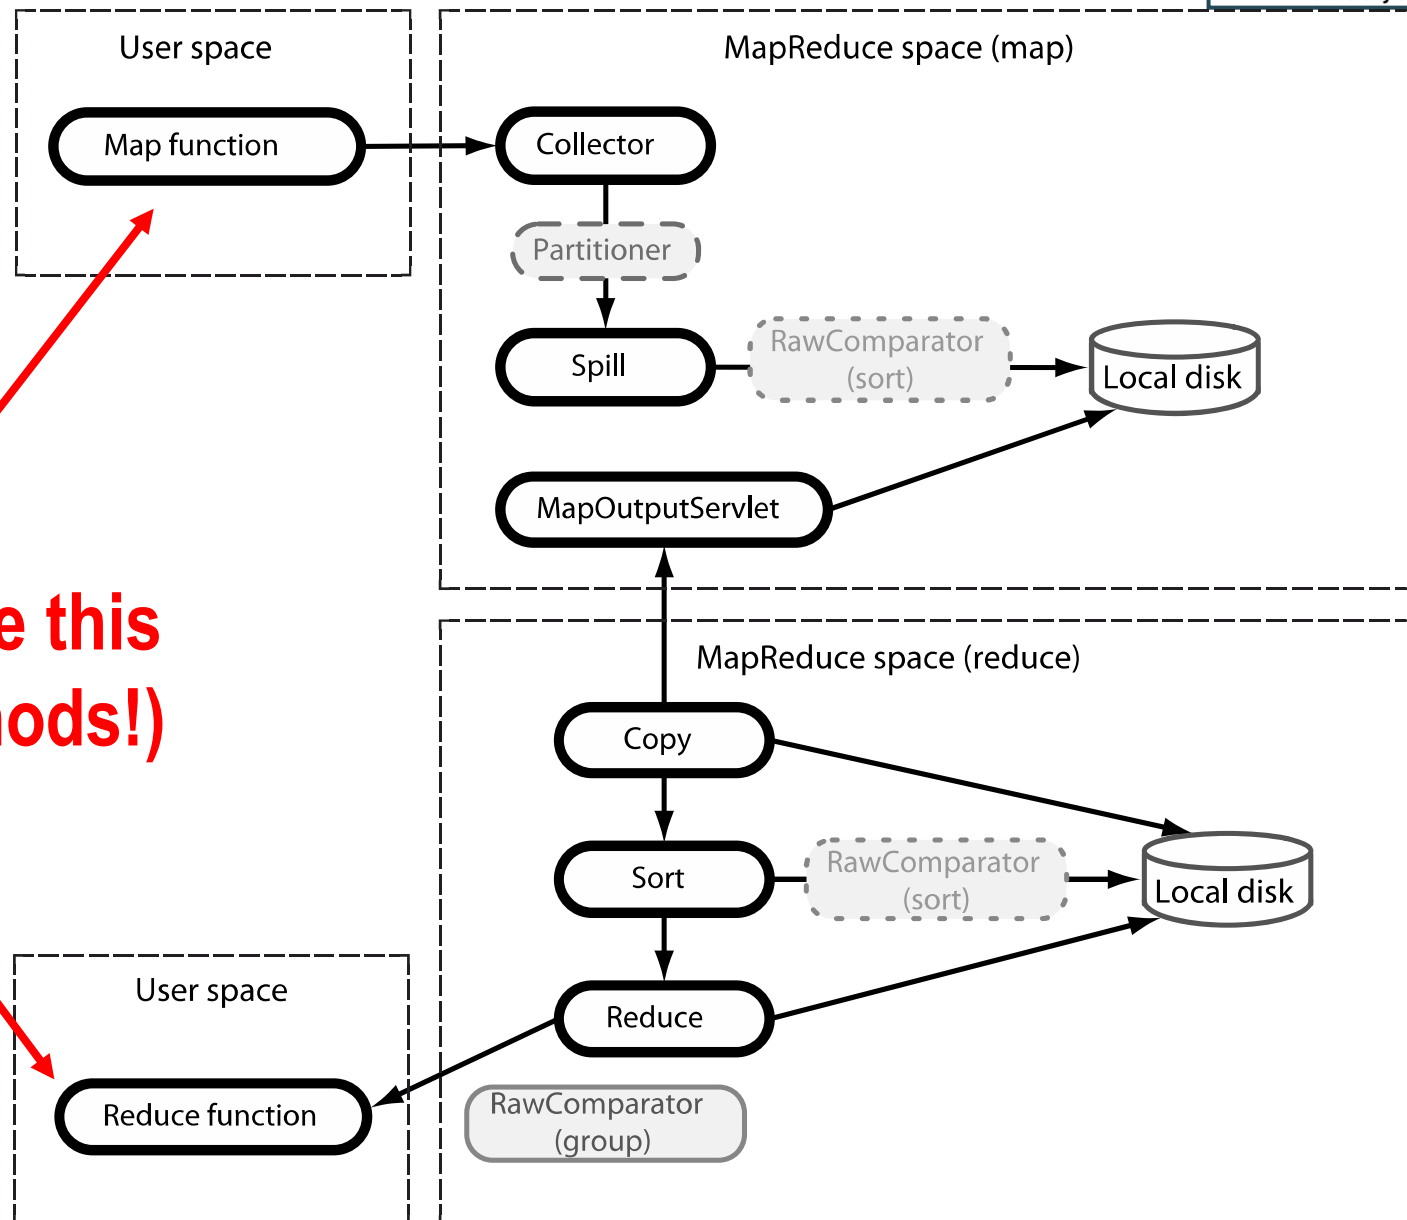
1. Split (logically) input data into computing chunks
2. Assign one chunk to a (co-located) NodeManager
3. Run 1..* **Mappers**
4. Shuffle and Sort
5. Run 1..* **Reducers**
6. Results from Reducers create the job output

Shuffle and Sort

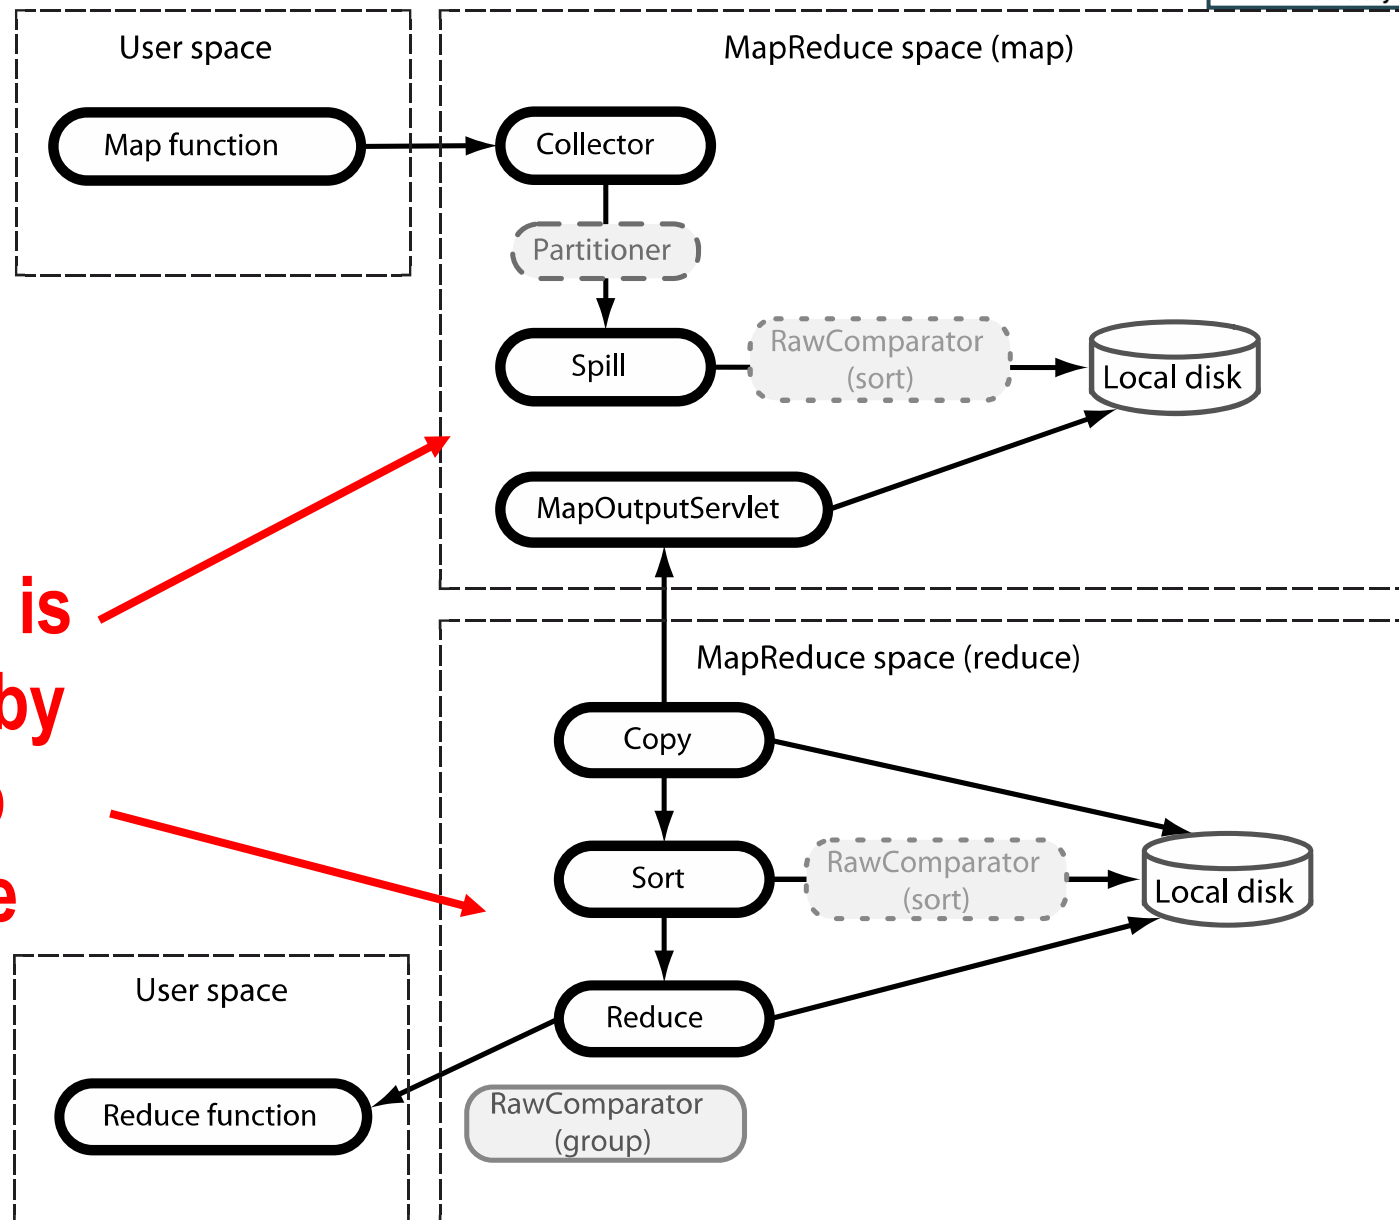
But we need nodes to exchange data...

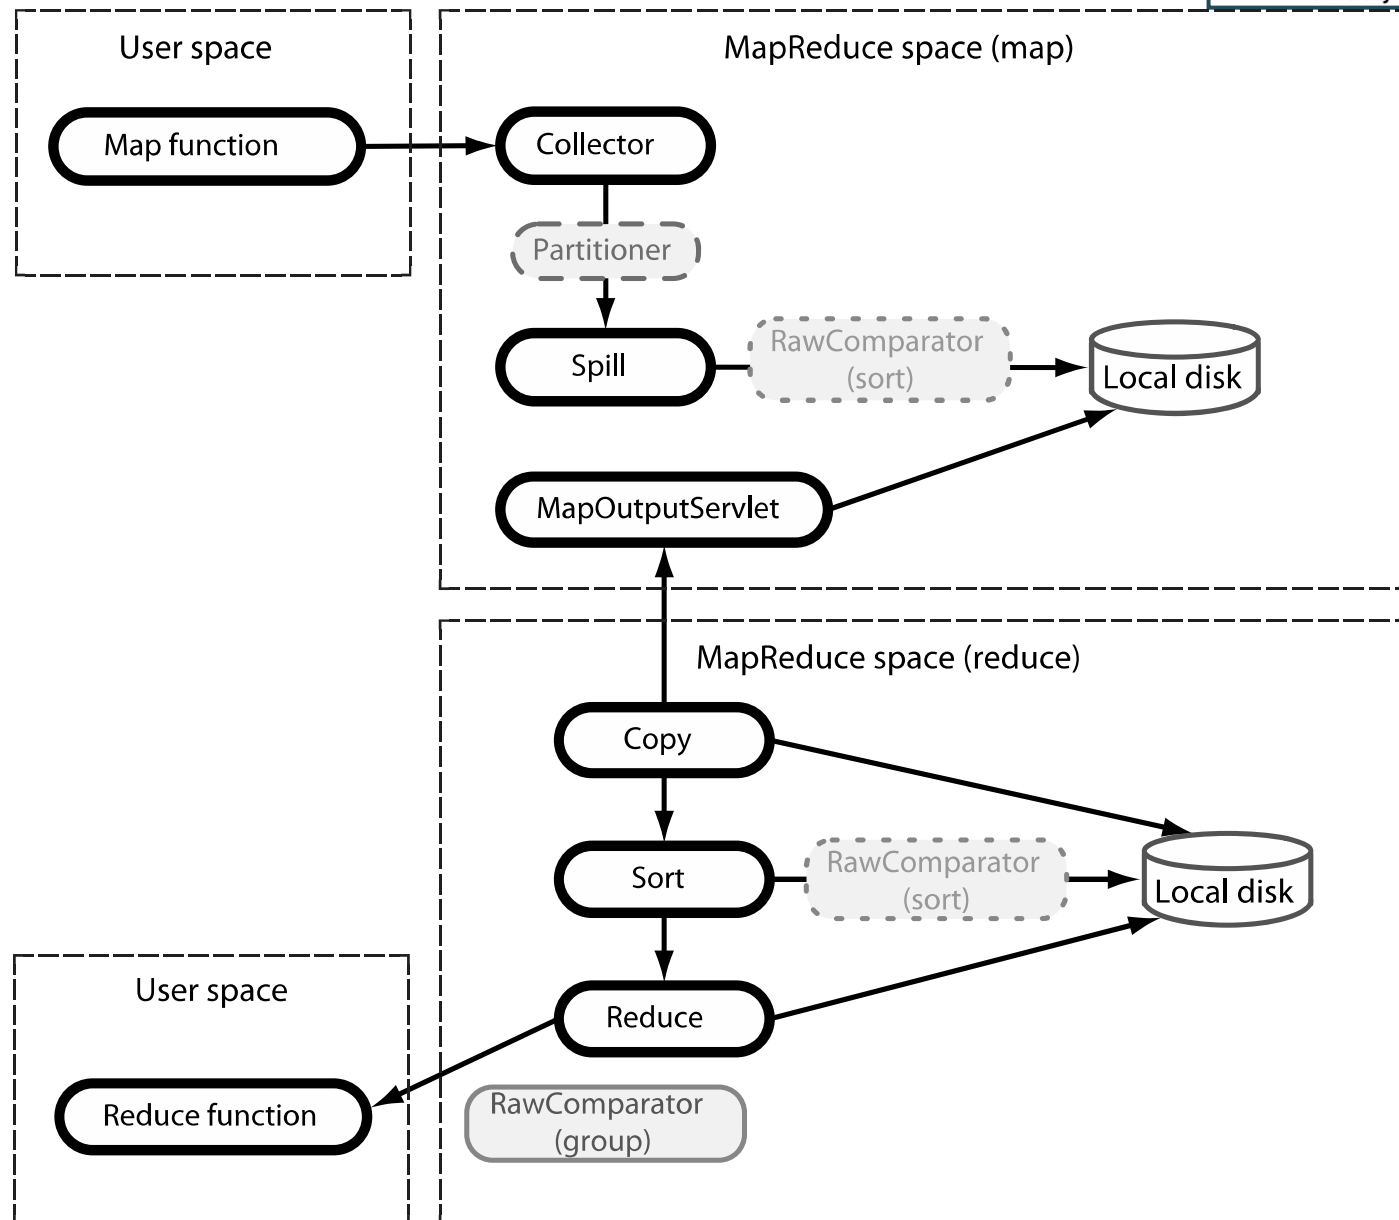


**We only write this
code (2 methods!)**



**This code is
included by
Hadoop
software**





Let's talk about a few things...

- Shuffling
 - Moving data from mappers to reducers
- Sorting
 - Ordering outputs before being processed by reducer

Shuffling: @Mapper

1. All key value pairs are **collected**
In-memory buffer (100MB default size), spills to HD
2. Pairs are **partitioned** depending on target reducer
Each partition is sorted by key
3. **Combiner** runs on each partition
4. Output is available to the Reducers through HTTP server threads

Sort: @Reducer

1. The reducer **copies** output from mappers
 - Asks ApplicationManager for map output locations
2. Downloaded output is **merged** and **sorted** into the full input for the Reducer
 - List of $\langle k_2, \text{list}\langle v_2 \rangle \rangle$, sorted by k_2

Contents

- Anatomy of a MapReduce job
- **The Combiner**
- Apache Hadoop
- Hadoop job execution: YARN
- Hadoop storage: HDFS

The cost of communications

- Parallelizing Map and Reduce jobs allow algorithms to scale close to linearly
- One potential bottleneck for MapReduce programs is the cost of Shuffle and Sort operations
 - Data has to be copied over network communications
 - All the keys emitted by the mappers
 - Sorting large amounts of elements can be costly
- **Combiner** is an additional optional step that is executed before these steps

The Combiner

- The **combiner** acts as a **preliminary reducer**
- It is executed **at** each mapper node just before sending all the <key, value> pairs for shuffling
- Reduces the number of emitted items
 - Improves efficiency
- It *cannot* be mandatory (the algorithm must work correctly if the Combiner is not invoked)

Word count combiner

```
public void Combine (String key,  
                     List<Integer> values) {  
    int sum = 0;  
    for (Integer count: values) {  
        sum += count;  
    }  
    emit (key, sum) ;  
}
```

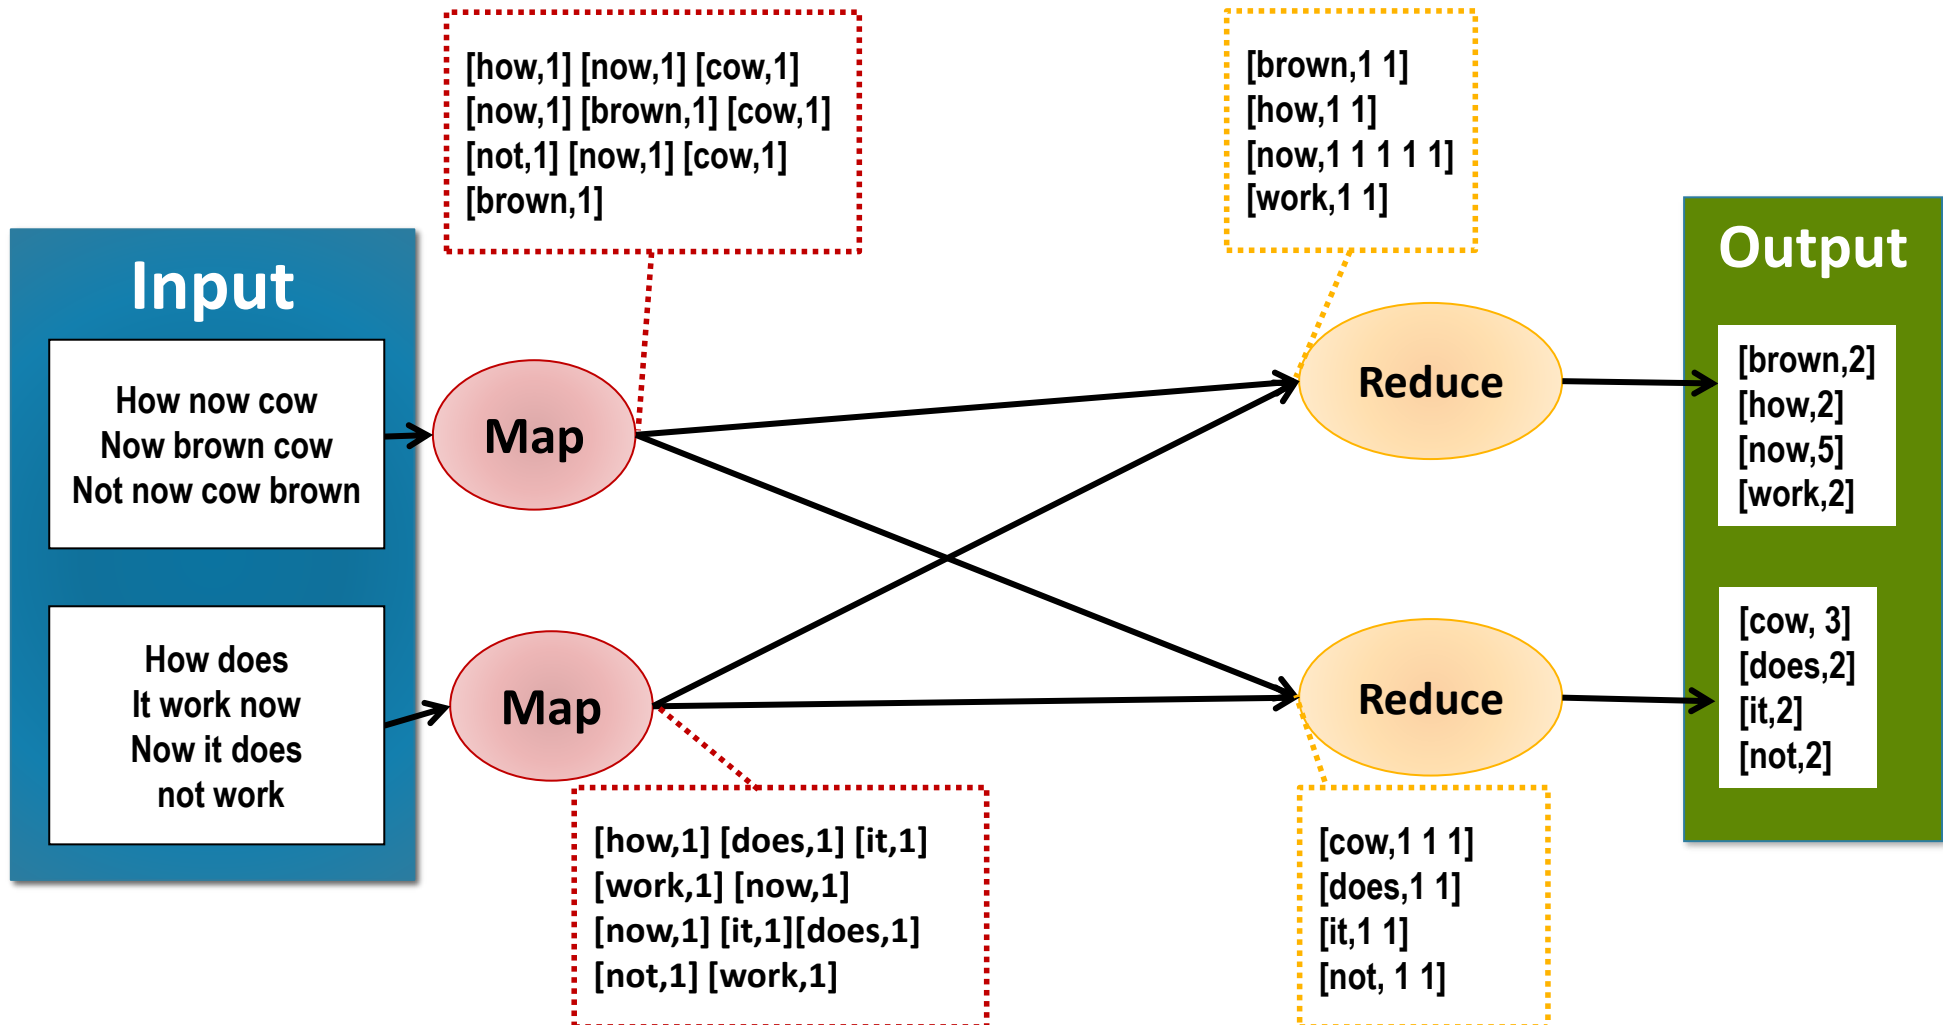
Remind you of anything?

It's the same as yesterday's reducer code!

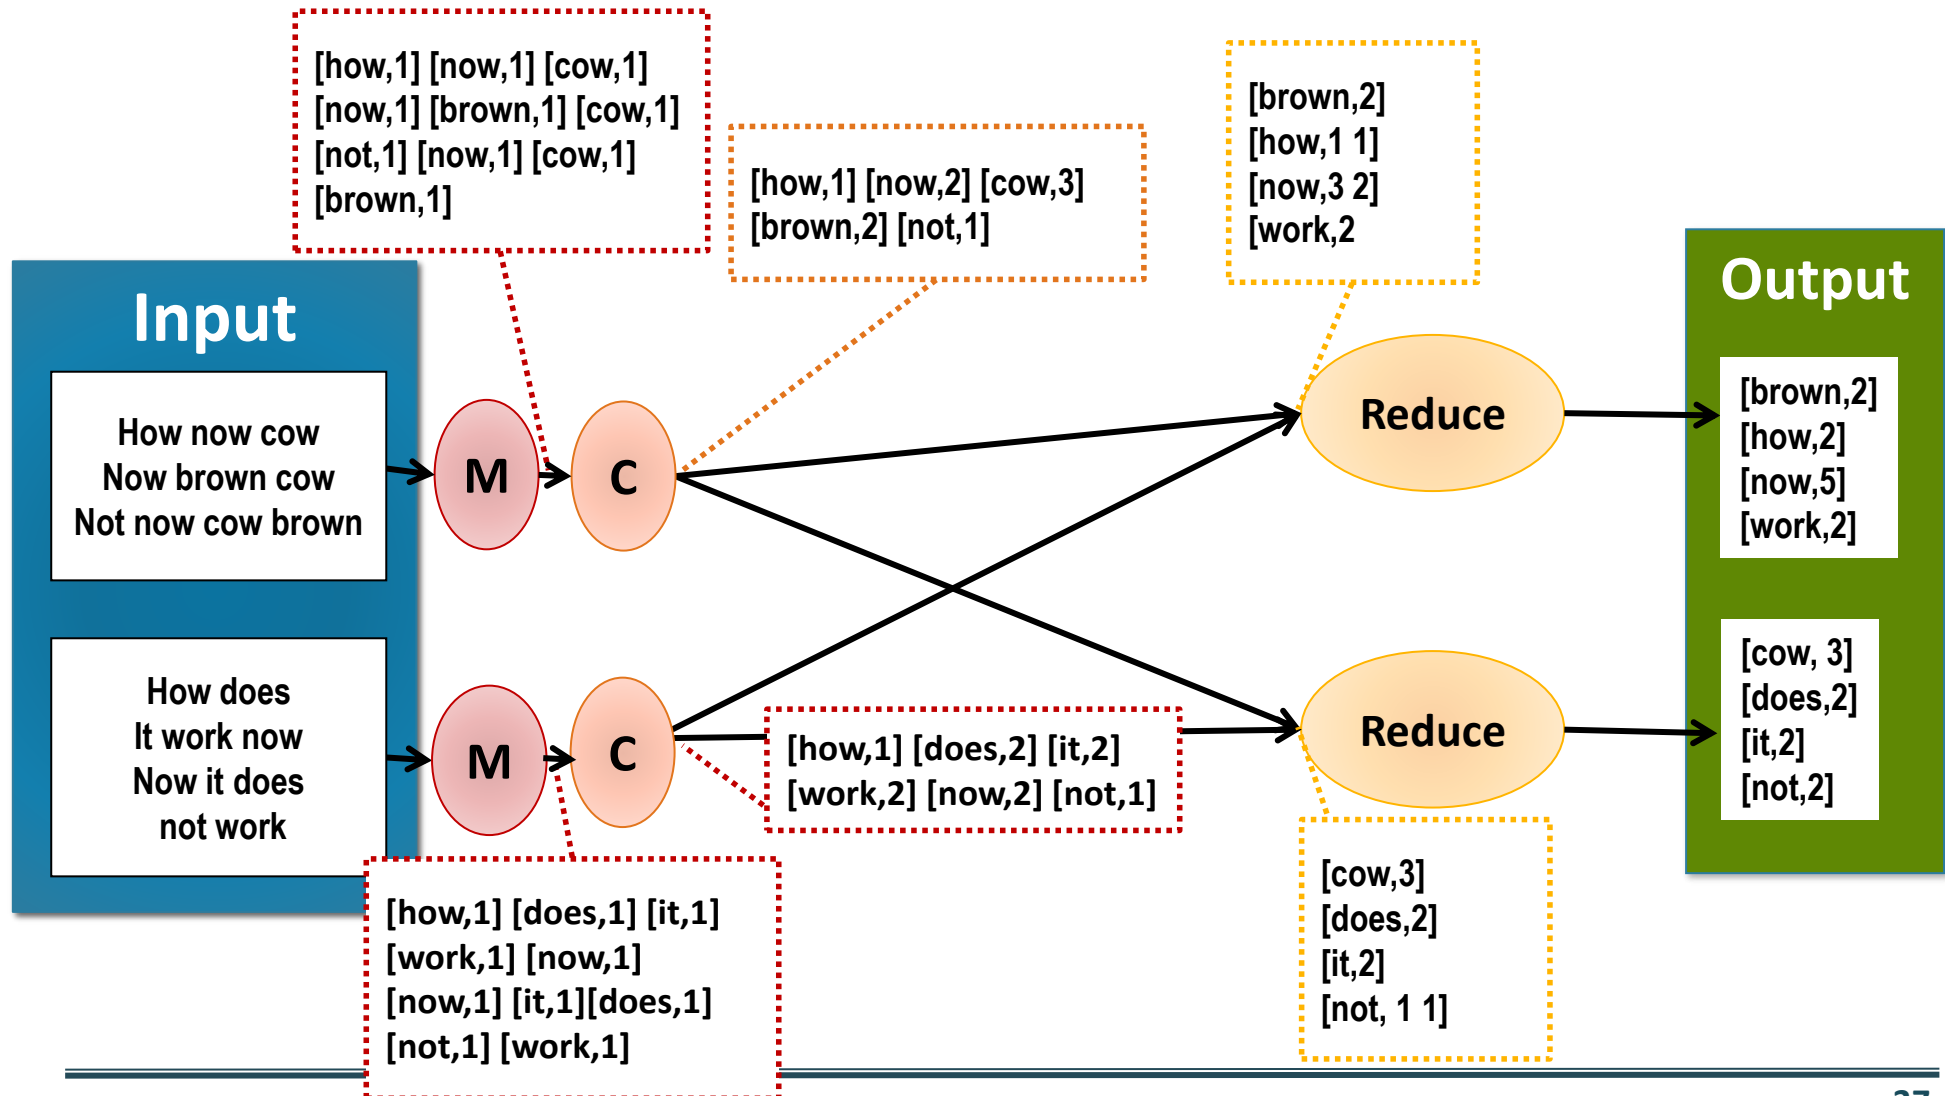
Combiner rules

- The combiner has the same structure as the reducer (same method signature) but must comply with these rules
 1. **Idempotent** - The number of times the combiner is applied can't change the output
 2. **Transitive** - The order of the inputs can't change the output
 3. **Side-effect free** - Combiners can't have side effects (or they won't be idempotent).
 4. **Preserve the sort order** - They can't change the keys to disrupt the sort order
 5. **Preserve the partitioning** - They can't change the keys to change the partitioning to the Reducers

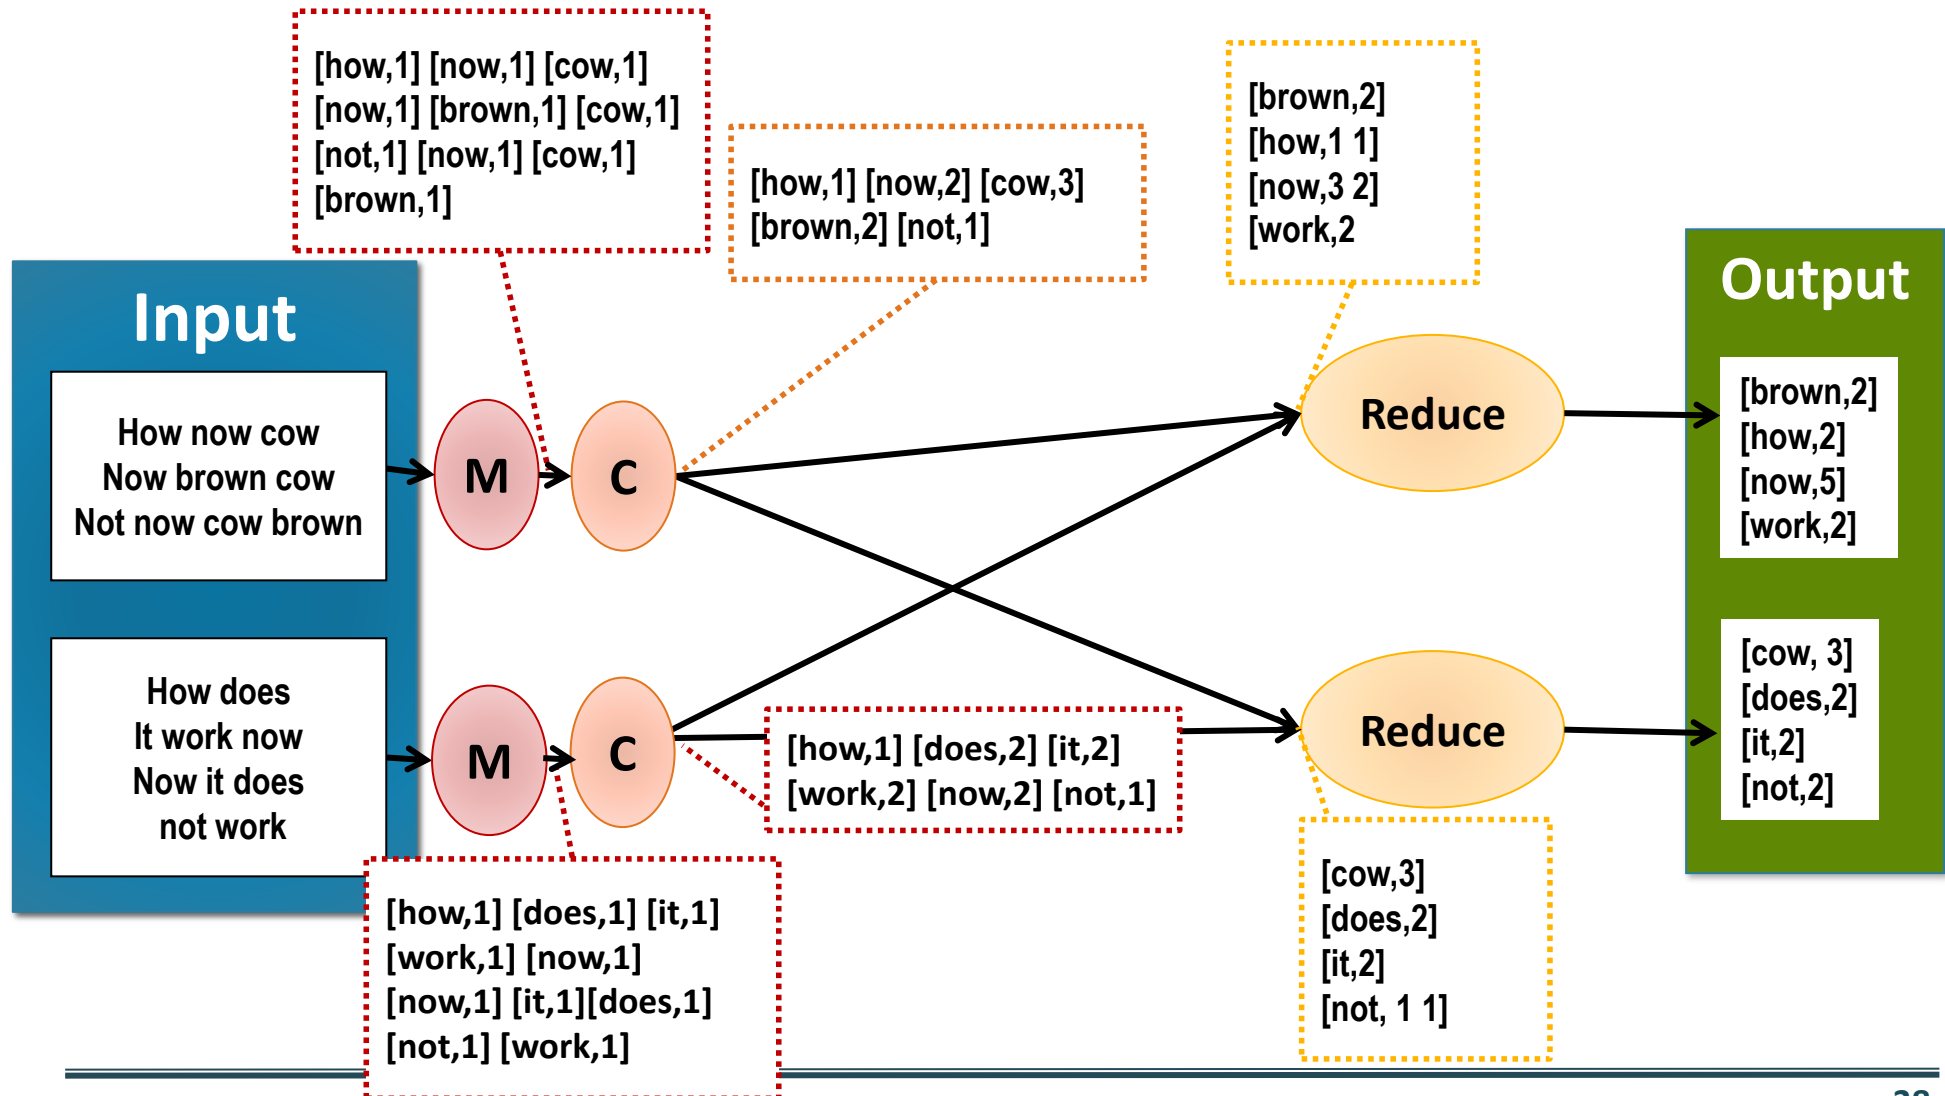
Word Count Example



Word Count Example with Combiner



Word Count Example with Combiner



Contents

- Anatomy of a MapReduce job
- The Combiner
- **Apache Hadoop**
- Hadoop job execution: YARN
- Hadoop storage: HDFS

The Apache Hadoop project

- Open source project hosted at Apache
- Implements the Map Reduce concept
- Written in Java
- Cloudera QuickStart VM includes it



Hadoop Architecture

- Hadoop executes on a cluster of networked PCs
- Each node runs a set of daemons

- ResourceManager
- NodeManager

Computing

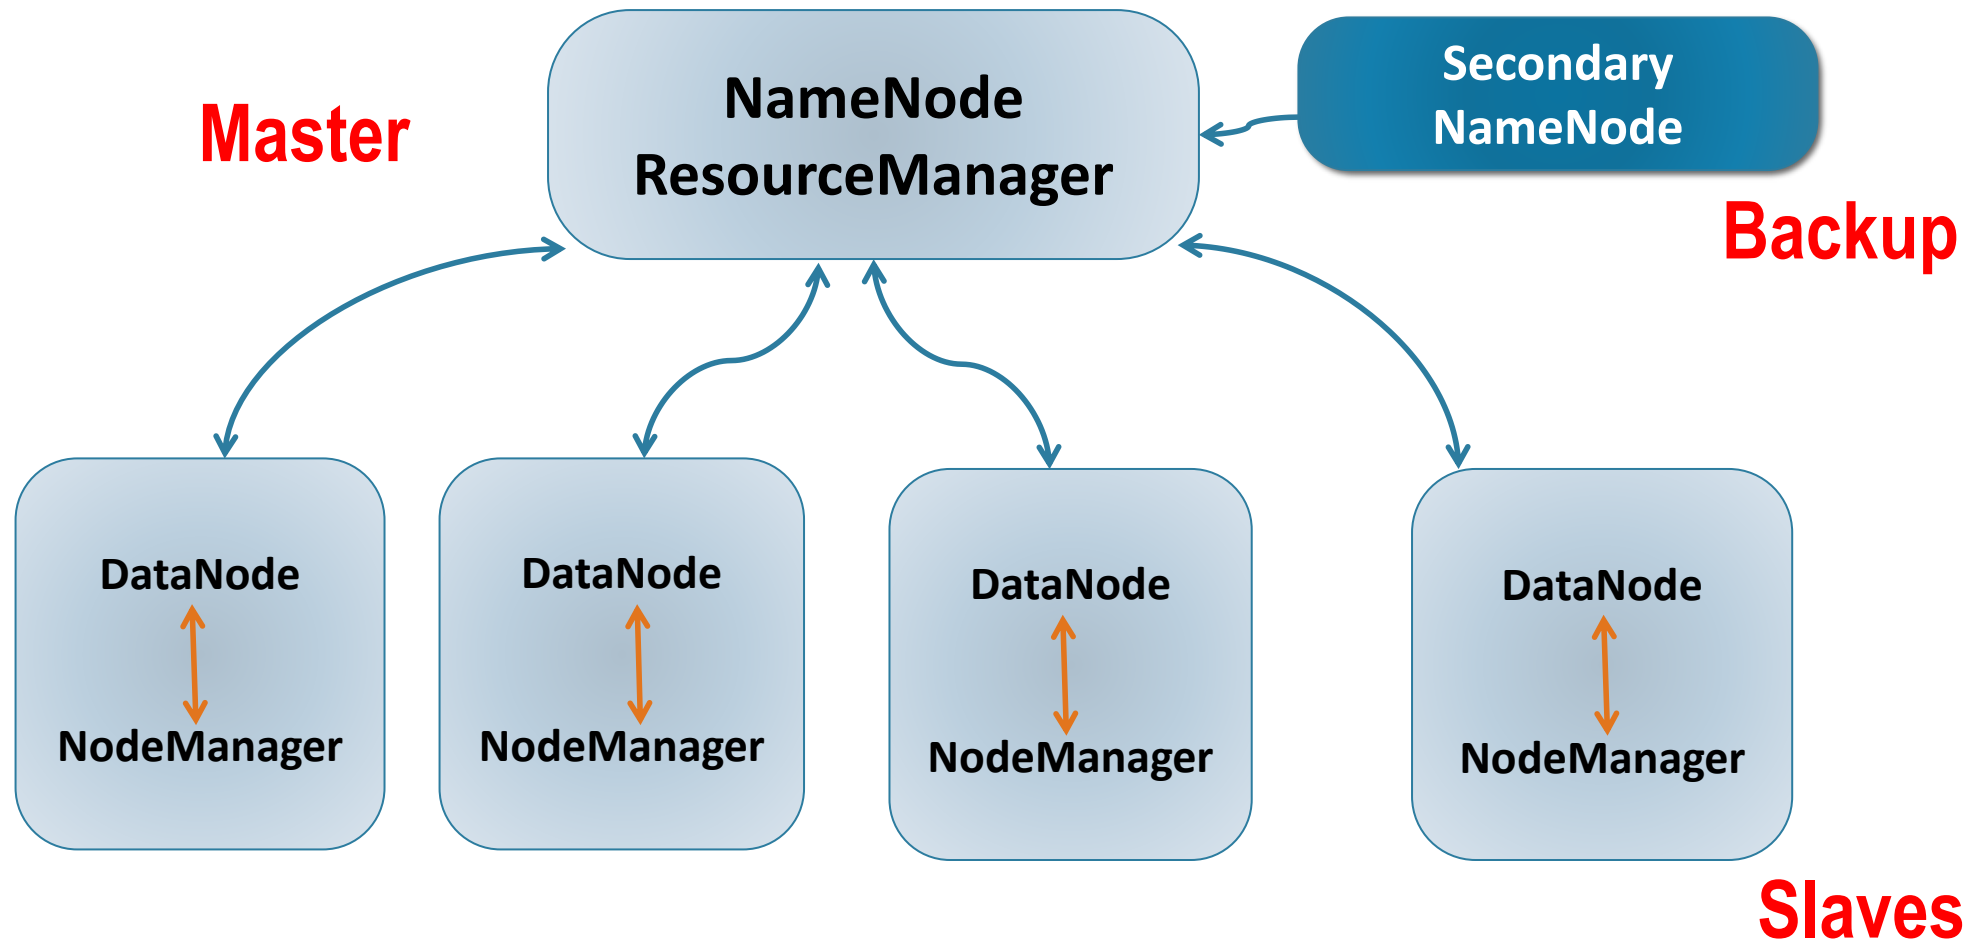
- NameNode
- SecondaryNameNode
- DataNode

Storage

Master Slave Architecture

- Master (just 1)
 - Is aware of all the slave nodes
 - Receives external requests
 - Decides who executes what, and when
 - Speaks with slaves
- Slave (1..*)
 - Worker node
 - Executes the tasks the master tells it to do

Hadoop Master-Slave architecture



Contents

- Anatomy of a MapReduce job
- The Combiner
- Apache Hadoop
- **Hadoop job execution: YARN**
- Hadoop storage: HDFS

Hadoop computation tasks

- Resource management
 - Being aware of what resources are in the cluster
 - Which resources are available now
- Job allocation
 - How many resources are needed to compute the job
 - Which nodes should execute each of the tasks
- Job execution
 - Coordinate task execution from workers
 - Make sure the job completes, deal with failures

Hadoop Job allocation

- Resource management needs to estimate **how many Map and Reduce tasks** are needed for a given job
 - Based on input dataset
 - Based on job definition
- Ideally, one different node will be allocated for each different Map/Reduce tasks
 - Otherwise multiple tasks can go to same node

Job execution: Complete MapReduce job flow

1. Split input data into computing chunks
2. Assign one chunk to a (co-located) NodeManager
3. Run 1..* **Mappers**
4. Shuffle and Sort
5. Run 1..* **Reducers**
6. Results from Reducers create the job output

How many Mappers are needed?

- Mapper parallelization:
 - Each Mapper processes a different **input split**
 - Input dataset size is known
- Number of mappers = input size / split size
 - If input has multiple small files, more Mappers can be invoked (Hadoop inefficiency)
 - If input size is 100MB and split size is 10MB...how many mappers?

How many Reducers are needed?

- Reducer parallelization
 - **Keys are partitioned** across the reducers
 - Hard to automatically estimate what is the right number
 - Too many Reducers can complicate too much shuffle and sort.
- Number of reducers = **User defined parameter**
 - Remember this for your Lab session!
 - It is in MapReduce job definition class

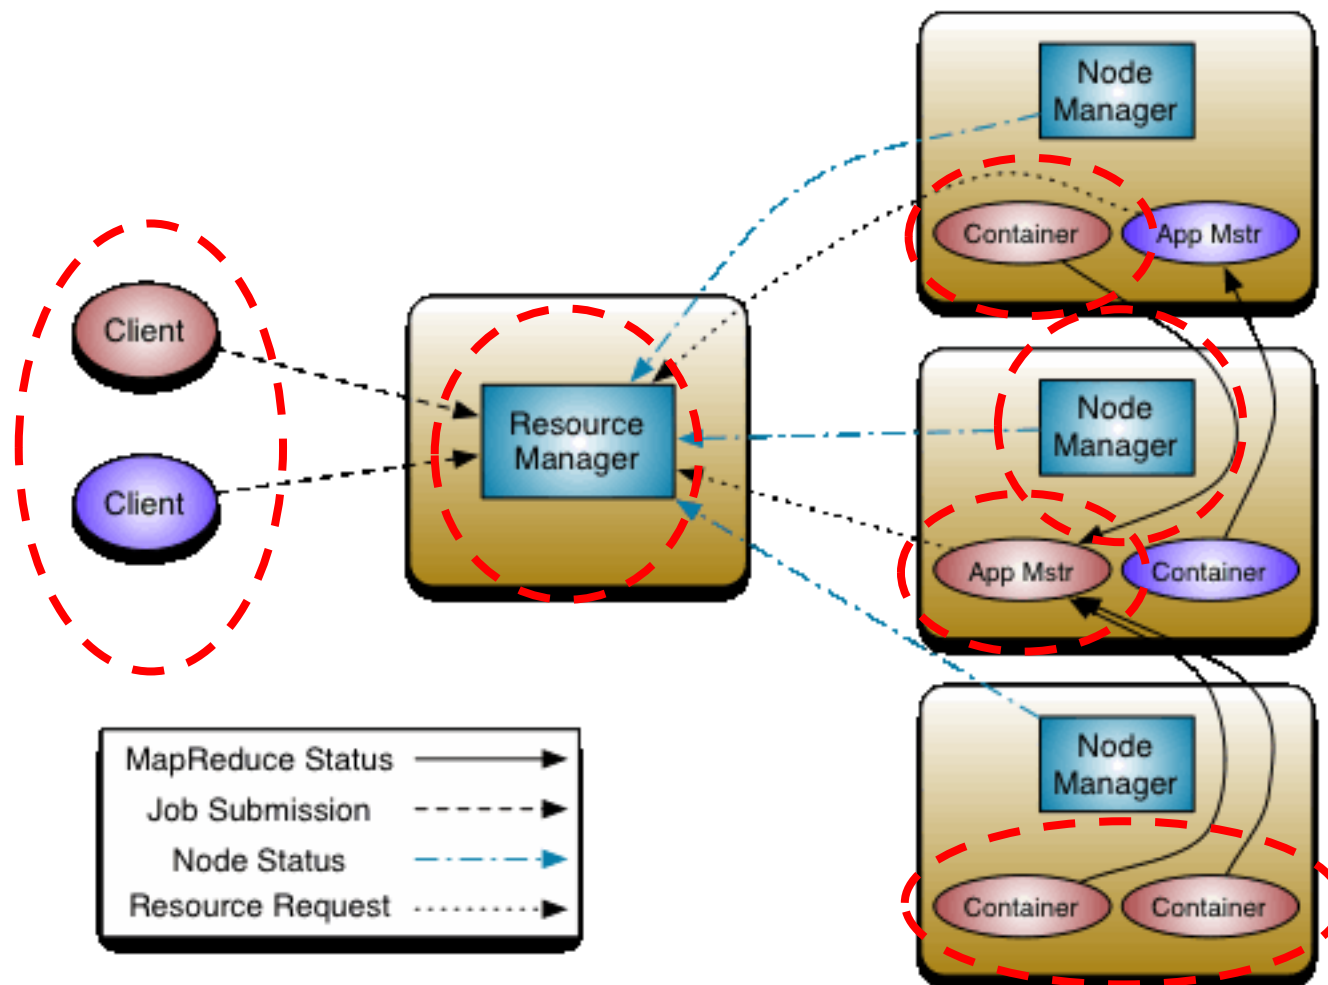
Hadoop execution daemons

- **ResourceManager** (1 per cluster)
 - Receives job requests from Hadoop Clients
 - Creates one **ApplicationMaster** per job to manage it
 - Allocates Containers in slave nodes, with assigned resources
 - Keeps track of health of NodeManager nodes
- **NodeManager** (1..* per cluster)
 - Coordinates execution of Map and Reduce tasks at node
 - Sends heartbeat messages to ResourceManager

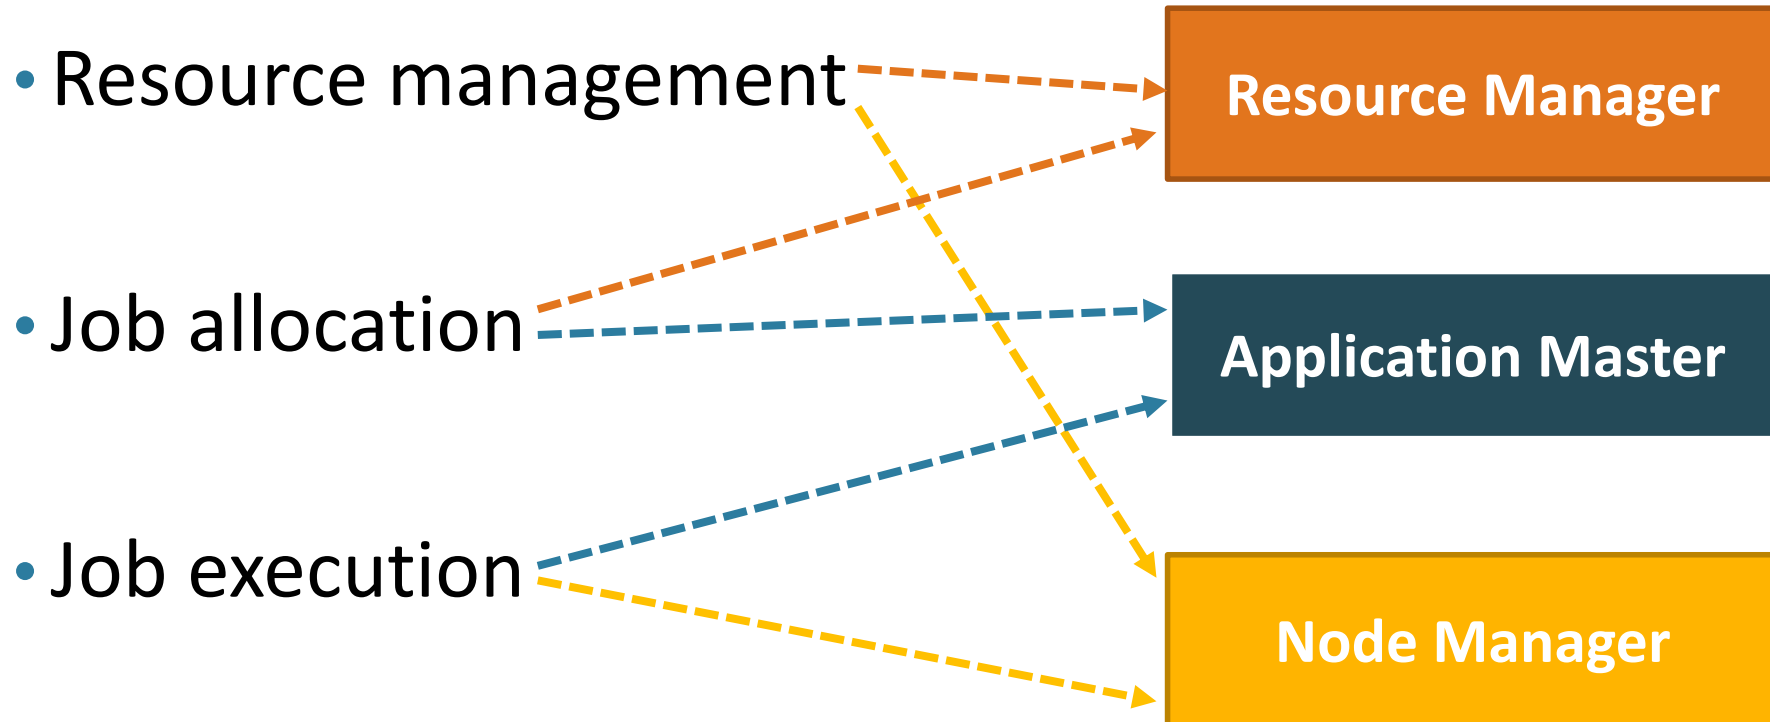
ApplicationMaster

- There is one ApplicationMaster per job.
- Implements the specific computing framework (e.g., MapReduce)
 - After creation, negotiates with ResourceManager how many resources will be required for the job
 - Decides which nodes will run Map and Reduce jobs among the Containers given by the ResourceManager
 - Reports to the ResourceManager about the progress and completion of the whole job
 - Is destroyed when the job is completed

Job Execution Architecture (YARN)



Responsibilities on computation tasks



Contents

- Anatomy of a MapReduce job
- The Combiner
- Apache Hadoop
- Hadoop job execution: YARN
- **Hadoop storage: HDFS**

HDFS

- HaDoo Distributed Filesystem
 - Shared storage among the nodes of the Hadoop cluster
- Storage for input and output of MapReduce jobs
- HDFS is Tailored for MapReduce jobs
 - Large block size (64MB default)
 - But not too large, blocks define the minimum parallelization unit
 - HDFS is not a POSIX compliant filesystem
 - Tradeoffs for improving data processing throughput

HDFS Data distribution

- Data distribution is a key element of the MapReduce model and architecture
- “Move computation to data” principle
 - Rather than copying data to the nodes that will process it, the jar is copied to the nodes where the data is stored
- Blocks are replicated over the cluster
 - Default ratio is three times
 - Spread replicas among different physical locations
 - Improves reliability

Hadoop Storage Daemons


- DataNode (1..* per cluster)
 - Stores blocks from the HDFS
 - Report periodically to NameNode list of stored blocks
- NameNode (1 per cluster)
 - Keeps index table with (all) the locations of each block
 - Heavy task, no computation responsibilities
 - Single point of failure
- Secondary Namenode (1 per cluster)
 - Communicates periodically with NameNode
 - Stores backup copy of index table

Data replication

Block Replication

Namenode (Filename, numReplicas, block-ids, ...)
/users/sameerp/data/part-0, r:2, {1,3}, ...
/users/sameerp/data/part-1, r:3, {2,4,5}, ...

Datanodes




Data replication

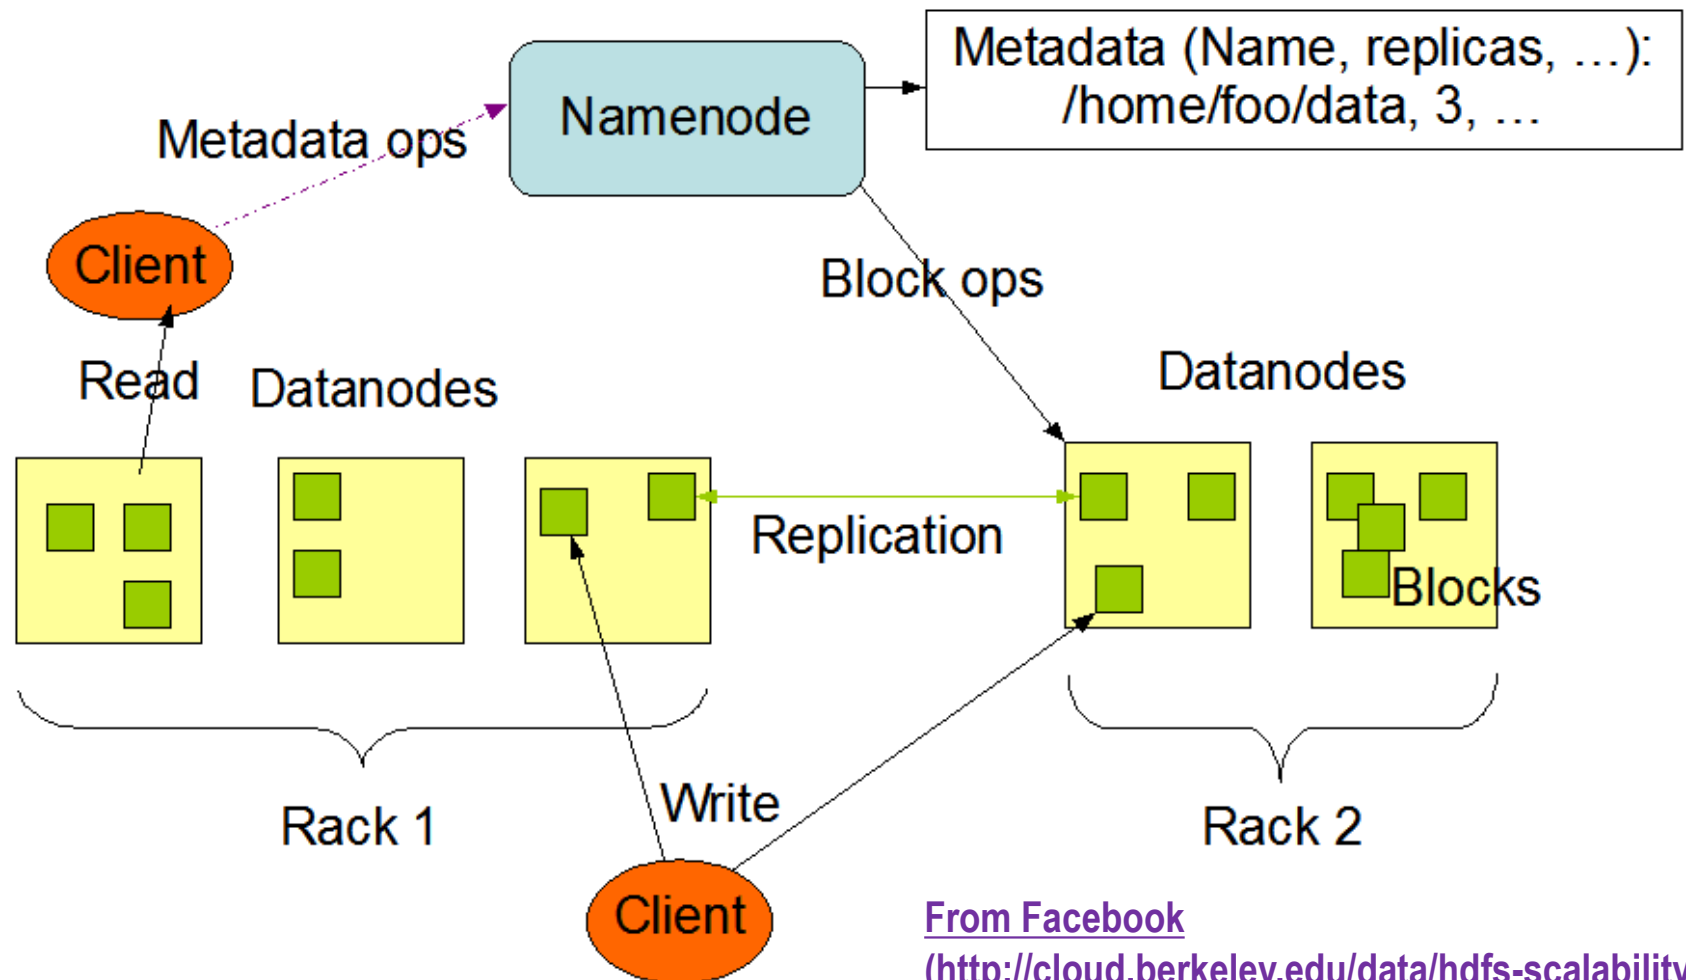
Block Replication

Namenode (Filename, numReplicas, block-ids, ...)
/users/sameerp/data/part-0, r:2, {1,3}, ...
/users/sameerp/data/part-1, r:3, {2,4,5}, ...

Datanodes

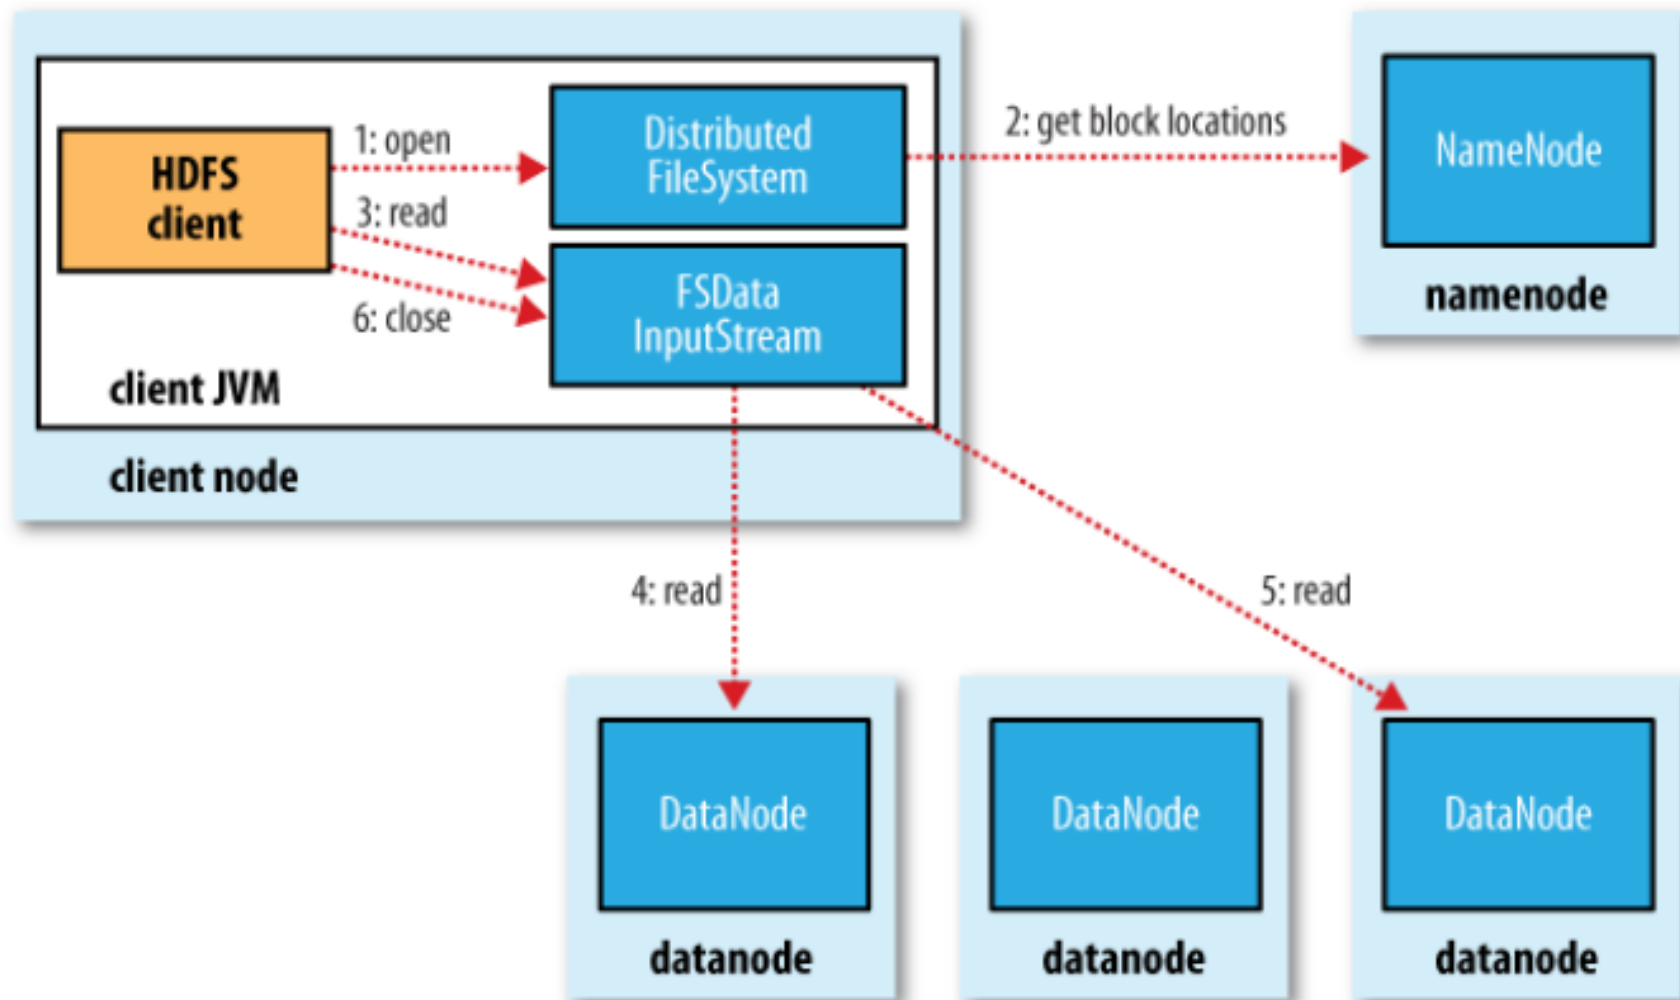


HDFS Usage

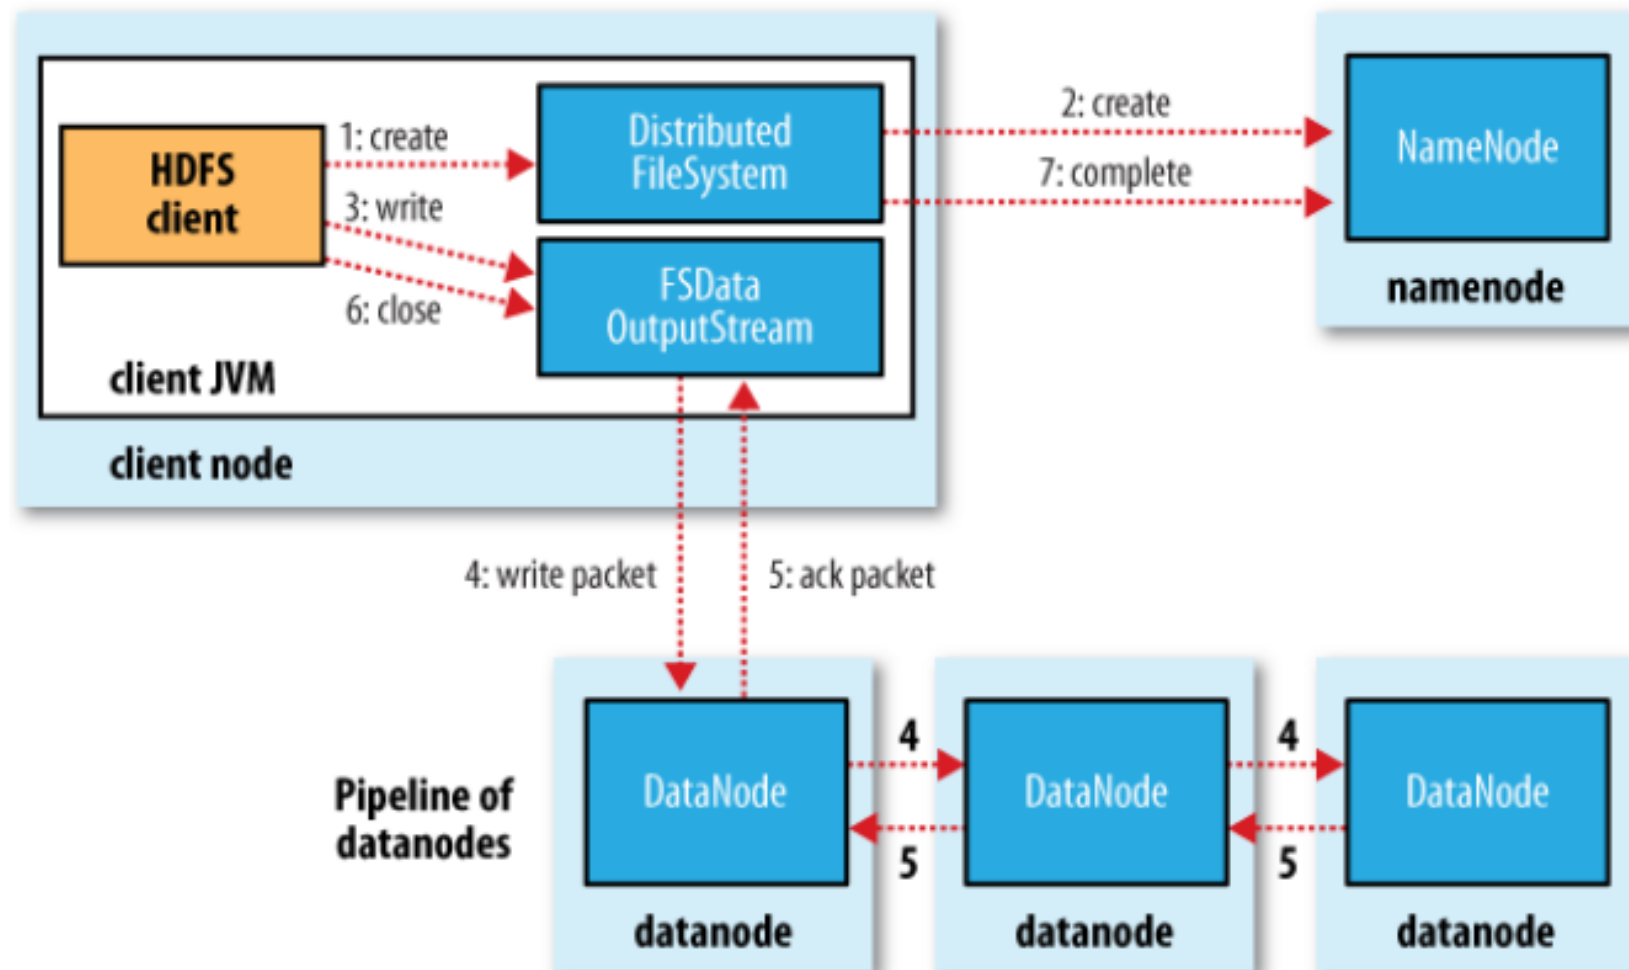


From Facebook
(<http://cloud.berkeley.edu/data/hdfs-scalability.pdf>)

HDFS File Read operation



HDFS File Write Operation



MR Job input and output data

1. Input data -> Mappers

- Mappers are assigned input splits from HDFS input path
 - (default 64MB)
- Data locality optimization: ApplicationManager attempts to assign Mappers where data block is stored

2. Reducers -> Output data

- Reducer output copied to HDFS
 - One file per Reducer
- For reliability concerns, HDFS replication

Recommended reading

- Hadoop YARN: Yet Another Resource Negotiator
 - <http://www.socc2013.org/home/program/a5-vavilapalli.pdf>
- How MapReduce works, Chapter 6, Hadoop: The Definitive Guide, 3rd Edition.
 - Available in QMUL Safaribooksonline
- HDFS design:
http://hadoop.apache.org/docs/r1.2.1/hdfs_design.html

Summary

- Anatomy of a MapReduce job
- The Combiner
- Apache Hadoop
- Hadoop job execution: YARN
- Hadoop storage: HDFS